

# File I/O

---

Prof. Joonwon Lee ([joonwon@skku.edu](mailto:joonwon@skku.edu))

TA – Jaehyung Park ([jaeseanpark@gmail.com](mailto:jaeseanpark@gmail.com))

TA – Luke Albano ([lukealbano@arcs.skku.edu](mailto:lukealbano@arcs.skku.edu))

Sungkyunkwan University

# Contents

- File/directory concept
- Unix I/O
- Standard I/O
- Error handling for file I/O

# File/Directory Concept



# What is File?

- A Linux **file** is a sequence of  $m$  bytes:
  - $B_0, B_1, \dots, B_k, \dots, B_{m-1}$
- All I/O devices are represented as files:
  - **/dev/sda1** (hard disk partition)
  - **/dev/tty2** (terminal)
    - Ctrl + Alt + F1 ~ F7
- Even the kernel is represented as a file:
  - **/boot/vmlinux-5.4.0-42-generic** (kernel image)
  - **/dev/mem** (kernel memory image)
  - **/proc** (kernel data structures)

# File Types

- Each file has a *type* indicating its role in the system
  - **Regular file**: contains arbitrary data
  - **Directory**: index for a related group of files
  - **Socket**: for communicating with a process on another machine
- Other file types beyond our scope
  - Named pipes (FIFOs)
  - Symbolic links
  - Character and block devices

# Regular File

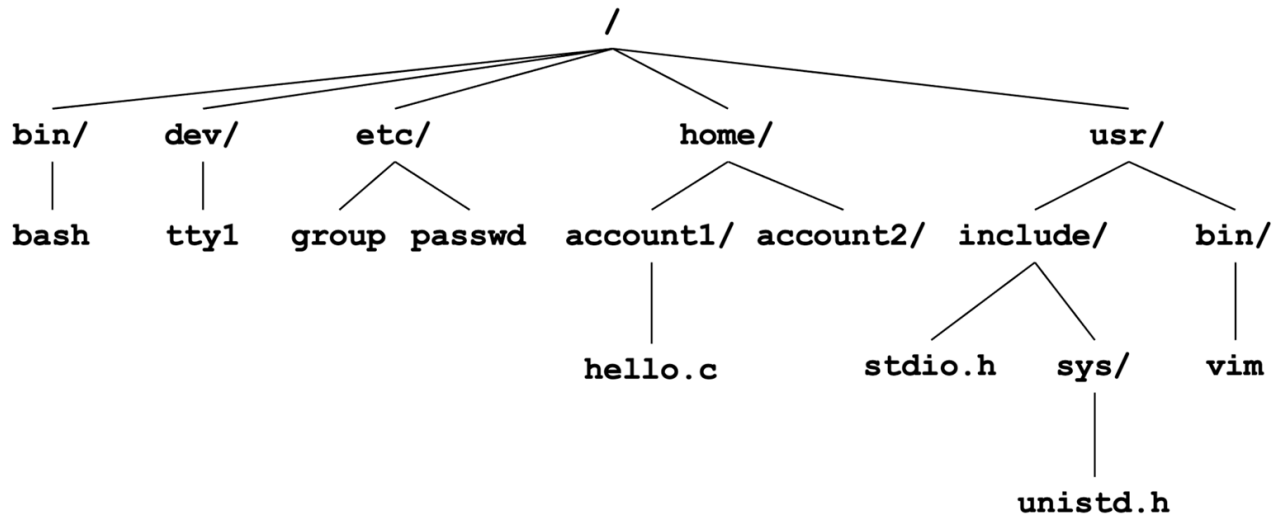
- A regular file contains arbitrary data
- **Applications** often distinguish between *text files* and *binary files*
  - Text files are regular files with only ASCII or Unicode characters
  - Binary files are everything else
    - e.g. object files, JPEG images
  - **Kernel doesn't know the difference!**
- Text file is sequence of text lines
  - Text line is sequence of *chars* terminated by newline *char* ('**\n**')

# Directory

- Directory consists of an array of *links*
  - Each link maps a *filename* to a file
- Each directory contains at least two entries
  - . (dot) is a link to itself
  - .. (dot dot, double dot) is a link to the parent directory in the directory hierarchy
- Commands for manipulating directories
  - **mkdir**: create empty directory
  - **ls** : view directory contents
  - **rmdir** : delete empty directory

# Directory Hierarchy

- All files are organized as a hierarchy anchored by root directory named '/' (slash)

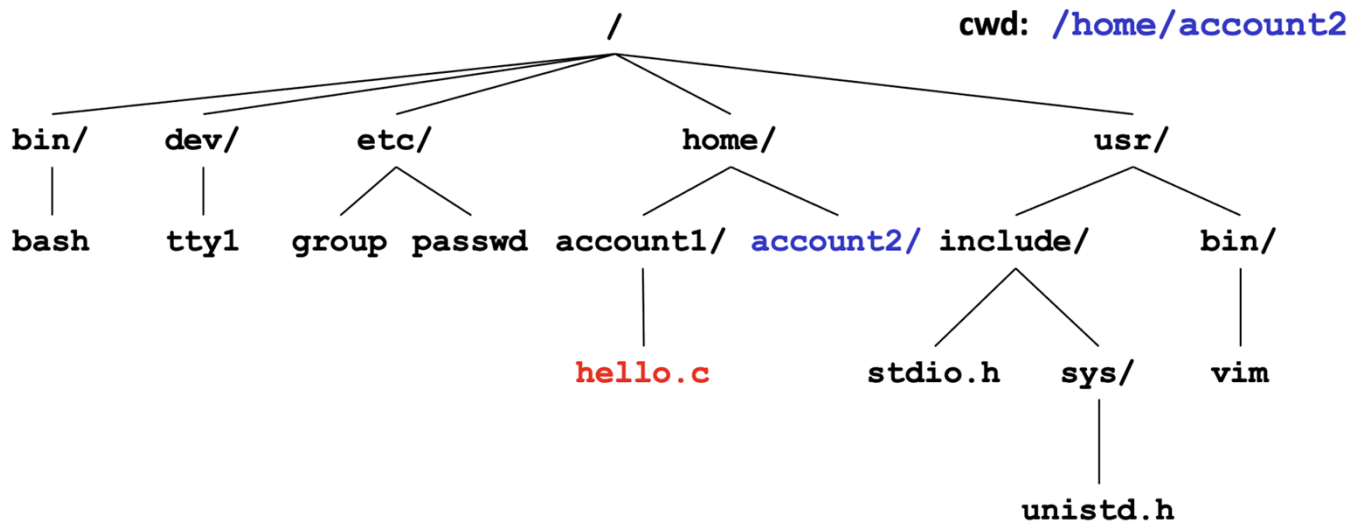


- Kernel maintains *current working directory* (cwd) for each process
  - Modified using the **cd** command



# Pathnames

- Locations of files in the hierarchy denoted by *pathnames*
  - Absolute pathname* starts with '/' and denotes path from root
    - `/home/account1/hello.c`
  - Relative pathname* denotes path from current working directory
    - `../home/account1/hello.c`

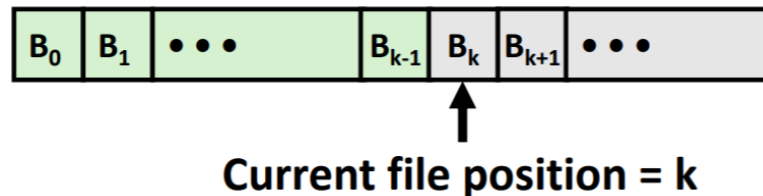




# Unix I/O

# Unix I/O Overview

- Kernel exports a simple interface called *Unix I/O*
  - Opening and closing files
    - open()** and **close()**
  - Reading and writing a file
    - read()** and **write()**
  - Changing the **current file position** (seek)
    - Indicates next offset into file to read or write
    - lseek()**



- All input and output is handled in a consistent and uniform way ("**byte stream**")

# Opening Files

- Opening a file informs the kernel that you are getting ready to access that file

```
int fd;    /* file descriptor */
if ((fd = open("/etc/hosts", O_RDONLY)) < 0) {
    perror("open");
    exit(1);
}
```

- Returns a small identifying integer *file descriptor*
  - `fd == -1` indicates that an error occurred
- Each process created by a Unix shell begins life with three open files associated with a terminal:
  - 0: standard input (stdin)
  - 1: standard output (stdout)
  - 2: standard error (stderr)

# Opening File Flags

- Flags for `open()` function
  - `O_RDONLY`: open for reading file
  - `O_WRONLY`: open for writing file
  - `O_RDWR`: open for reading and writing
  - `O_CREAT`: create the file if it doesn't exist
  - `O_APPEND`: append to the end of file
- Flags when create file with `O_CREAT` flag
  - **mode** flag locates third argument position in open system call
    - `int open(const char *pathname, int flags, mode_t mode)`
  - e.g. `S_IRWXU`, `S_IRUSR`, `S_IRGRP`, `S_IROTH`
    - 0764 equals `S_IRWXU` | `S_IRGRP` | `S_IWGRP` | `S_IROTH`

# Closing Files

- Closing a file informs the kernel that you are finished accessing that file

```
int fd;      /* file descriptor */
int retval; /* return value */

if ((retval = close(fd)) < 0) {
    perror("close");
    exit(1);
}
```

- Closing an already closed file is a recipe for disaster in threaded programs (more on this later)
- Moral: **Always check return codes**, even for seemingly benign functions such as `close()`

# Reading Files

- Reading a file copies bytes from the *current file position* to *memory*, and then updates file position

```
char buf[512];
int fd;      /* file descriptor */
int nbytes;  /* number of bytes read */

/* Open file fd ... */
/* Then read up to 512 bytes from file fd */
if ((nbytes = read(fd, buf, sizeof(buf))) < 0) {
    perror("read");
    exit(1);
}
```

- Returns number of bytes read from file fd into buf
  - **nbytes < 0** indicates that an error occurred
  - **Short counts** (**nbytes < sizeof(buf)**) are possible and are not errors!

# Writing Files

- Writing a file copies bytes from *memory* to the *current file position*, and then updates current file position

```
char buf[512];
int fd;      /* file descriptor */
int nbytes;  /* number of bytes read */

/* Open the file fd ... */
/* Then write up to 512 bytes from buf to file fd */
if ((nbytes = write(fd, buf, sizeof(buf)) < 0) {
    perror("write");
    exit(1);
}
```

- Returns number of bytes written from `buf` to file `fd`
  - **`nbytes < 0`** indicates that an error occurred
  - As with reads, **short counts** are possible and are not errors!



# File Offset

- An offset of an opened file can be set explicitly by calling **lseek()**, **lseek64()**

```
char buf[512];
int fd;      /* file descriptor */
off_t pos;   /* file offset */

/* Get current file offset */
pos = lseek(fd, 0, SEEK_CUR);
/* The file offset is incremented by written bytes */
write(fd, buf, sizeof(buf));
/* Set file position to the first byte of the file */
pos = lseek(fd, 0, SEEK_SET);
```

- Returns the new offset of the file fd
  - **nbytes < 0** indicates that an error occurred
  - An offset can be set beyond the end of the file
    - If data is written at that point, a file "hole" is created

# Simple Unix I/O Example

- Copying standard input to standard output one byte at a time

```
int main(void) {  
    char c;  
    while(read(STDIN_FILENO, &c, 1) != 0)  
        write(STDOUT_FILENO, &c, 1);  
    exit(0);  
}
```

# Dealing with Short Counts

- Short counts can occur in these situations:
  - Encountering (end-of-file) EOF on reads
  - Reading text lines from a terminal
  - Reading and writing network sockets
- Short counts does not occur in these situations:
  - Reading from disk files (except for EOF)
  - Writing to disk files
- Best practice is to always allow for short counts

# Handling Short Counts

```
ssize_t rio_readn(int fd, void *usrbuf, size_t n) {
    size_t nleft = n;
    ssize_t nread;
    char *bufp = usrbuf

    while(nleft > 0)
        if ((nread = read(fd, bufp, nleft)) < 0) {
            if (errno == EINTR) /* interrupted by signal handler return */
                nread = 0; /* and call read() again */
            else
                return -1; /* errno set by read() */
        }
        else if (nread == 0)
            break; /* EOF */
        nleft -= nread;
        bufp += nread;
    }
    return (n - nleft); /* return >= 0 */
}
```

# Accessing Directory

- Only recommended operation on a directory: read its **entries**
  - **dirent** structure contains information about a directory entry

```
struct dirent {  
    ino_t      d_ino;      /* inode number */  
    off_t      d_off;      /* offset to the next dirent */  
    unsigned short d_reclen; /* length of this record */  
    unsigned char d_type;   /* type of file */  
    char        d_name[256]; /* filename */  
};
```

- **DIR** structure contains information about directory while stepping through its entries (see example in the next slide)
- **opendir()** and **closedir()**
  - Open and close directory stream
- **readdir()**
  - Return a pointer to a **dirent** structure representing directory entry
  - Return **NULL** on reaching the end of the directory stream

# Example of Accessing Directory

```
#include <sys/types.h>
#include <dirent.h>
#include <stdio.h>
#include <stdlib.h>

/* dircheck.c - Opening directory and reading its entries */
int main(int argc, char **argv) {
    DIR *directory;
    struct dirent *de;

    if (!(directory = opendir(argv[1]))) {
        perror("Failed to open directory");
        exit(1);
    }

    while (0 != (de = readdir(directory)))
        printf("Found: %s\n", de->d_name);

    closedir(directory);
    exit(0);
}
```

# File Metadata

- **Metadata** is data about data, in this case file data
  - Per-file metadata maintained by kernel, accessed by users with the **stat** and **fstat** functions

```
/* Metadata returned by the stat and fstat functions */
struct stat {
    dev_t      st_dev;      /* device */
    ino_t      st_ino;      /* inode */
    mode_t     st_mode;     /* protection and file type */
    nlink_t    st_nlink;    /* number of hard links */
    uid_t      st_uid;      /* user ID of owner */
    gid_t      st_gid;      /* group ID of owner */
    dev_t      st_rdev;     /* device type (if inode device) */
    off_t      st_size;     /* total size, in bytes */
    unsigned long st_blksize; /* blocksize for filesystem I/O */
    unsigned long st_blocks; /* number of blocks allocated */
    time_t     st_atime;     /* time of last file access */
    time_t     st_mtime;     /* time of last file modification */
    time_t     st_ctime;     /* time of last inode change */
};                          /* statbuf.h included by sys/stat.h */
```

# System Calls for File Metadata

- An information about a file can be obtained explicitly by calling **stat()**, **fstat()**, **lstat()**
  - Three header files must be included

```
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>
```
  - **int stat(const char \*path, struct stat \*buf)**
    - Fills file metadata in buf pointed by path
  - **int fstat(int fd, struct stat \*buf)**
    - Identical to **stat()**, except that the file is specified by the file descriptor fd
  - **int lstat(const char \*path, struct stat \*buf)**
    - Identical to **stat()**, except that if path is a symbolic link, fill link itself, not the file



# Example of Accessing File Metadata

```
/* statcheck.c - Querying and manipulating a file's meta data */
```

```
int main (int argc, char **argv) {
```

```
    struct stat st;
```

```
    char *type, *readok;
```

```
    stat(argv[1], &st);
```

```
    /* Determine file type */
```

```
    if (S_ISREG(st.st_mode)) type = "regular";
```

```
    else if (S_ISDIR(st.st_mode)) type = "directory";
```

```
    else type = "other";
```

```
    /* Check read access */
```

```
    if ((st.st_mode & S_IRUSR)) readok = "yes";
```

```
    else readok = "no";
```

```
    printf("type: %s, read: %s\n", type, readok);
```

```
    exit(0);
```

```
}
```

```
linux> ./statcheck statcheck.c
type: regular, read: yes
linux> chmod 000 statcheck.c
linux> ./statcheck statcheck.c
type: regular, read: no
```



# Standard I/O

# Standard I/O Functions

- The C standard library (`libc.so`) contains a collection of higher-level **standard I/O** functions
- Examples of standard I/O functions:
  - Opening and closing files (**`fopen()`** and **`fclose()`**)
  - Reading and writing bytes (**`fread()`** and **`fwrite()`**)
  - Reading and writing text lines (**`fgets()`** and **`fputs()`**)
  - Formatted reading and writing (**`fscanf()`** and **`fprintf()`**)

# Standard I/O Streams

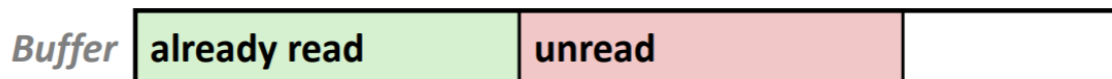
- Standard I/O models open files as **streams**
  - Abstraction for a file descriptor and a buffer in memory
- C programs begin life with three open streams (defined in **stdio.h**)
  - **stdin** (standard input)
  - **stdout** (standard output)
  - **stderr** (standard error)

```
#include <stdio.h>
extern FILE *stdin; /* standard input (descriptor 0) */
extern FILE *stdout; /* standard output (descriptor 1) */
extern FILE *stderr; /* standard error (descriptor 2) */

int main() {
    fprintf(stdout, "Hello, world\n");
}
```

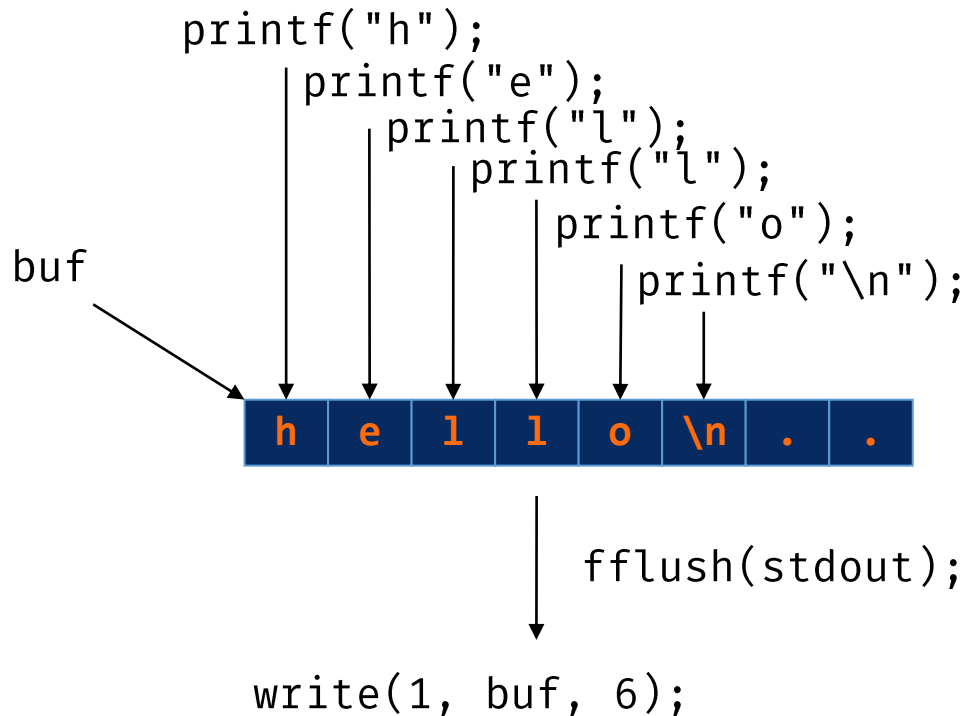
# Buffered I/O : Motivation

- Applications often read/write one character or line at a time
  - **getc()**, **putc()**, **ungetc()**
  - **gets()**, **fgets()**
    - Read line of text one character at a time, stopping at newline
- Implementing as Unix I/O calls expensive
  - Each **read()** and **write()** require Unix kernel calls
    - > 10,000 clock cycles
- Solution : buffered read
  - Use Unix **read()** to grab block of bytes
  - User input functions take one byte at a time from buffer
    - Refill buffer when empty



# Buffering in Standard I/O

- Standard I/O functions use buffered I/O



- Buffer flushed to output fd on "\n", call to **fflush()** or **exit()**, or return from main

# Standard I/O Buffering in Action

- You can see this buffering in action for yourself, by using Linux program, **strace**.

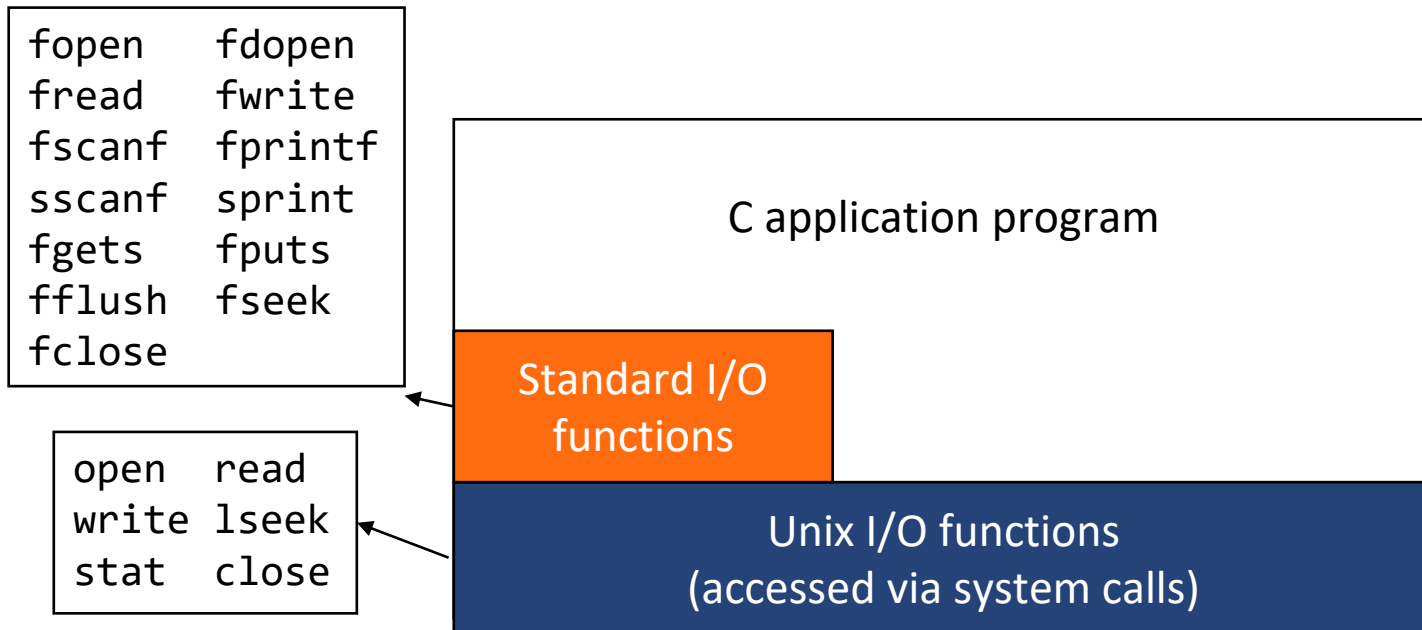
```
#include <stdio.h>

int main() {
    printf("h");
    printf("e");
    printf("l");
    printf("l");
    printf("o");
    printf("\n");
    fflush(stdout);
    exit(0);
}
```

```
linux> strace ./hello
execve("./hello", ["hello"], [/* ... */]).
...
write(1, "hello\n", 6)                = 6
...
exit_group(0)                         = ?
```

# Unix I/O vs. Standard I/O

- Standard I/O are implemented using low-level Unix I/O



- Which ones should you use in your programs?



# Pros and Cons of Unix I/O

## ■ Pros

- The most general and lowest overhead form of I/O
  - All other I/O packages are implemented using Unix I/O functions
- Unix I/O provides functions for accessing *file metadata*

## ■ Cons

- System call overheads for small-sized I/O
- Dealing with short counts is tricky and error prone
- Efficient reading of text lines requires some form of buffering, also tricky and error prone
- These issues are addressed by the standard I/O

# Pros and Cons of Standard I/O

## ■ Pros

- Buffering increases efficiency by decreasing the number of **read()** and **write()** system calls
- Shout counts are handled automatically

## ■ Cons

- Provides no function for accessing file metadata
- Standard I/O is not appropriate for input and output on network sockets
  - But there is a way using **fdopen()**



# Error Handling for File I/O

# *perror* Function / *errno* Variable

- When a system call fails,
  - Returns -1 (or NULL for certain library functions)
  - The latest error information is stored in "*errno*"
- Variable *errno*
  - Stores int value indicating cause of error
  - int type extern global variable
  - Thread-safe!
- Function *perror*
  - Explain error information stored in *errno*
  - Print out the information through stderr stream.

# errno.h

`/usr/include/asm-generic/errno-base.h`

File: `/usr/include/asm-generic/errno-base.h`

```
1  /* SPDX-License-Identifier: GPL-2.0 WITH Linux-syscall-note */
2  #ifndef _ASM_GENERIC_ERRNO_BASE_H
3  #define _ASM_GENERIC_ERRNO_BASE_H
4
5  #define EPERM      1 /* Operation not permitted */
6  #define ENOENT     2 /* No such file or directory */
7  #define ESRCH     3 /* No such process */
8  #define EINTR     4 /* Interrupted system call */
9  #define EIO       5 /* I/O error */
10 #define ENXIO     6 /* No such device or address */
11 #define E2BIG     7 /* Argument list too long */
12 #define ENOEXEC   8 /* Exec format error */
13 #define EBADF     9 /* Bad file number */
14 #define ECHILD   10 /* No child processes */
15 #define EAGAIN   11 /* Try again */
16 #define ENOMEM   12 /* Out of memory */
17 #define EACCES   13 /* Permission denied */
18 #define EFAULT   14 /* Bad address */
19 #define ENOTBLK  15 /* Block device required */
20 #define EBUSY    16 /* Device or resource busy */
21 #define EEXIST    17 /* File exists */
22 #define EXDEV    18 /* Cross-device link */
23 #define ENODEV   19 /* No such device */
24 #define ENOTDIR  20 /* Not a directory */
25 #define EISDIR   21 /* Is a directory */
26 #define EINVAL   22 /* Invalid argument */
27 #define ENFILE   23 /* File table overflow */
28 #define EMFILE   24 /* Too many open files */
29 #define ENOTTY   25 /* Not a typewriter */
30 #define ETXTBSY  26 /* Text file busy */
31 #define EFBIG    27 /* File too large */
32 #define ENOSPC   28 /* No space left on device */
33 #define EPIPE    29 /* Illegal seek */
34 #define EROFS    30 /* Read-only file system */
35 #define EMLINK   31 /* Too many links */
36 #define EPIPE    32 /* Broken pipe */
37 #define EDOM     33 /* Math argument out of domain of func */
38 #define ERANGE   34 /* Math result not representable */
39
40 #endif
```

- 1) Defined as integers
- 2) Stored in **errno** when error occurs
- 3) Accessible through "**man errno**"

# Handling Errors Example

```
#include <stdio.h>
#include <stdlib.h>
#include <errno.h>

int main() {
    FILE *fp;
    fp = fopen("file.txt", "r");
    if(fp == NULL){
        fprintf(stderr, "errno: %d\n", errno);
        perror("Error");
        exit(1);
    }
    fclose(fp);
    exit(0);
}
```

```
linux> ./test
linux> errno: 2
linux> Error: No such file or directory
```

# Summary

- Unix file I/O
  - **open()**, **read()**, **write()**, **close()**, ...
  - A uniform way to access files, I/O devices, network sockets, kernel data structures, Etc.
- When to use standard I/O
  - When working with disk or terminal files
- When to use raw Unix I/O
  - When you need to fetch file metadata
  - When you read or write network sockets or pipes
  - In rare cases when you need absolute highest performance

# Remind

- 6 System calls
  - `open()`
  - `close()`
  - `read()`
  - `write()`
  - `lseek()`
  - `stat()` / `fstat()`



# File I/O Example

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;

fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6);
read(fd, buffer+6, 2);
```

```
lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);
```

```
lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

**File state (FD: 3)**

**path: "hello-dos.txt"**  
**position: 0**  
**size: 20**



H	e	l	l	o	.	\r	\n
W	h	o		a	r	e	
y	o	u	?				

# File I/O Example

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;

fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6); // "Hello."
read(fd, buffer+6, 2);
```

```
lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);
```

```
lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

**File state (FD: 3)**

**path: "hello-dos.txt"**  
**position: 6**  
**size: 20**



H	e	l	l	o	.	\r	\n
W	h	o		a	r	e	
y	o	u	?				

# File I/O Example

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;

fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6);
read(fd, buffer+6, 2); // "Hello.\r\n"
```

File state (FD: 3)


path: "hello-dos.txt"  
position: 8  
size: 20

```
lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);
```

```
lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```



	e	l	l	o	.	\r	\n
W	h	o		a	r	e	
y	o	u	?				

# File I/O Example

```
char filename[] = "hello-dos.txt";  
int fd;  
char buffer[16];  
off_t pos = 0; // long long;
```

```
fd = open(filename, O_RDWR | O_CREAT, 0755);  
read(fd, buffer, 6);  
read(fd, buffer+6, 2);
```

```
lseek(fd, -2, SEEK_CUR);  
buffer[0] = '\n';  
write(fd, buffer, 1);
```

```
lseek(fd, 8, SEEK_SET);  
strcpy(buffer, "How");  
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);  
if (fd < 0)  
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

File state (FD: 3)

path: "hello-dos.txt"  
position: 6  
size: 20



H	e	l	l	o	.	\r	\n
W	h	o		a	r	e	
y	o	u	?				

# File I/O Example

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;
```

```
fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6);
read(fd, buffer+6, 2);
```

```
lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);
```

```
lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

File state (FD: 3)

path: "hello-dos.txt"  
position: 7  
size: 20



H	e	l	l	o	.	\n	\n
W	h	o		a	r	e	
y	o	u	?				

# File I/O Example

```
char filename[] = "hello-dos.txt";  
int fd;  
char buffer[16];  
off_t pos = 0; // long long;
```

```
fd = open(filename, O_RDWR | O_CREAT, 0755);  
read(fd, buffer, 6);  
read(fd, buffer+6, 2);
```

```
lseek(fd, -2, SEEK_CUR);  
buffer[0] = '\n';  
write(fd, buffer, 1);
```


```
lseek(fd, 8, SEEK_SET);  
strcpy(buffer, "How");  
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);  
if (fd < 0)  
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

File state (FD: 3)

path: "hello-dos.txt"  
position: 8  
size: 20



	e	l	l	o	.	\n	\n
W	h	o		a	r	e	
y	o	u	?				

# File I/O Example

```
char filename[] = "hello-dos.txt";  
int fd;  
char buffer[16];  
off_t pos = 0; // long long;
```

```
fd = open(filename, O_RDWR | O_CREAT, 0755);  
read(fd, buffer, 6);  
read(fd, buffer+6, 2);
```

```
lseek(fd, -2, SEEK_CUR);  
buffer[0] = '\n';  
write(fd, buffer, 1);
```

```
lseek(fd, 8, SEEK_SET);  
strcpy(buffer, "How");  
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);  
if (fd < 0)  
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

File state (FD: 3)

path: "hello-dos.txt"  
position: 11  
size: 20



H	e	l		o	.	\n	\n
W	h	o		a	r	e	
y	o	u	?				

# File I/O Example

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;

fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6);
read(fd, buffer+6, 2);

lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);

lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);

close(fd);

fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

**File state (FD: 3)  
:CLOSED**



# File I/O Example

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;

fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6);
read(fd, buffer+6, 2);

lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);

lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);

close(fd);

fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

# Lab Exercise

- Add line numbers
  - Create “p4” directory in your home directory
  - Go to the directory and download “Aladdin.txt”
    - wget <http://www.fpx.de/fp/Disney/Scripts/Aladdin.txt>
  - Create “Aladdin\_num.txt”, which has line numbers in front of each line of text
  - Only use **open()**, **read()**, **write()**
  - **Don't use other functions such as, scanf(), printf(), sprintf(), fscanf(), fprintf(), and fopen()**
    - **Otherwise, zero**

# Exercise

```
ALADDIN: THE COMPLETE SCRIPT
COMPILED BY BEN SCRIPPS <34RQNPQ@CMUVM.CSV.CMICH.EDU>
(Portions Copyright (c) 1992 The Walt Disney Company)

PEDDLER:  Oh I come from a land
          From a faraway place
          Where the caravan camels roam
          Where they cut off your ear /Where it's flat and immense
          If they don't like your face /And the heat is intense
          It's barbaric, but hey--it's home!
          When the wind's at your back
          And the sun's from the west
          And the sand in the glass is right
          Come on down,
          Stop on by
          Hop a carpet and fly
          To another Arabian night!
```

Aladdin.txt

```
1 | ALADDIN: THE COMPLETE SCRIPT
2 | COMPILED BY BEN SCRIPPS <34RQNPQ@CMUVM.CSV.CMICH.EDU>
3 | (Portions Copyright (c) 1992 The Walt Disney Company)
4 |
5 | PEDDLER:  Oh I come from a land
6 |          From a faraway place
7 |          Where the caravan camels roam
8 |          Where they cut off your ear /Where it's flat and immense
9 |          If they don't like your face /And the heat is intense
10 |         It's barbaric, but hey--it's home!
11 |         When the wind's at your back
12 |         And the sun's from the west
13 |         And the sand in the glass is right
14 |         Come on down,
15 |         Stop on by
16 |         Hop a carpet and fly
17 |         To another Arabian night!
```

Aladdin\_num.txt

```
linux> ls
ex4
linux> wget http://www.fpx.de/fp/Disney/Scripts/Aladdin.txt
linux> ls
Aladdin.txt ex4
linux> ./ex4 Aladdin.txt
linux> ls
Aladdin_num.txt Aladdin.txt ex4
```

# Exercise

- Submit your exercise source code and Makefile
  - InUiYeJi cluster
  - Submit the folder into p4
  - ~swe2024-41\_23s/bin/submit p4 p4
  - We will compile by using command *make*
    - If compilation fails, your points for this exercise will be zero

./p4

- Makefile
- main.c
- Other files you need to build executable
- **Remove “Aladdin.txt” and “Aladdin\_num.txt”**

# Summary Report

- Summary report about man command result of
  - **open()**
  - **read()**
  - **write()**
  - **close()**
- Submission form
  - A4 size PDF format (No page limitation)
  - [SWE2024 Report-3] studentID\_name
  - Ex) [SWE2024 Report-3] 2022XXXXXX\_홍길동
  - Submit via iCampus
  - Due: until Friday, 24 March 2023, 23:59