

Process

Prof. Joonwon Lee (joonwon@skku.edu)

TA – Jaehyung Park (jaeseanpark@gmail.com)

TA – Luke Albano (lukealbano@arcs.skku.edu)

Sungkyunkwan University

Process (1)

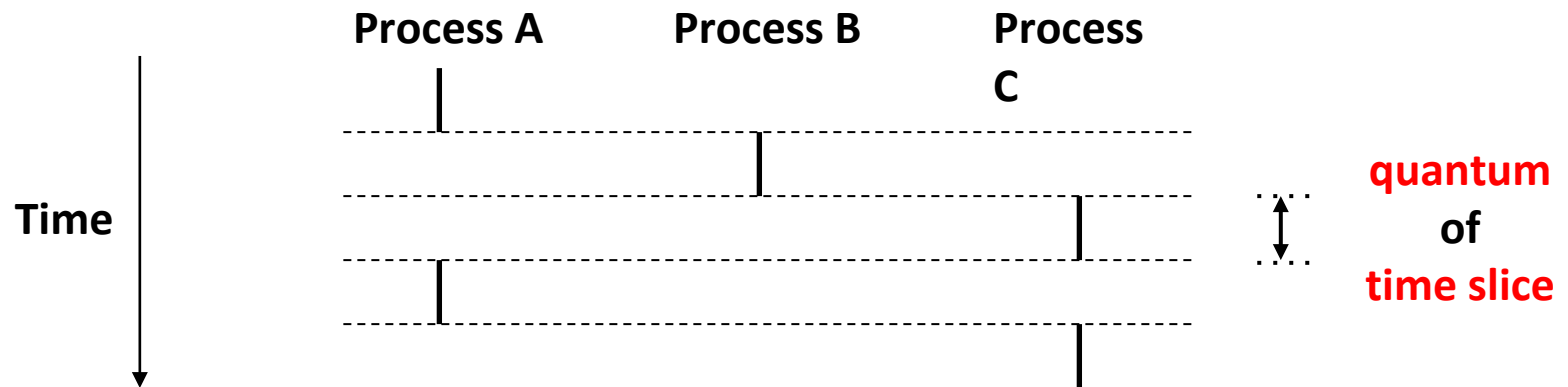
- It is not same as "program" or "processor"
- What's the difference between ...
 - Program
 - **Process**
 - Processor
- Process
 - **An instance of a program in execution**
 - One of the most profound ideas in computer science

Process (2)

- Process provides each program with two key abstractions:
 - Logical control flow
 - Each program seems to have exclusive use of the CPU
 - Private address space
 - Each program seems to have exclusive use of main memory
- How are these illusions maintained?
 - Process executions interleaved (multitasking)
 - Address space managed by virtual memory system


Logical Control Flows (1)

- Each process has its own logical control flow



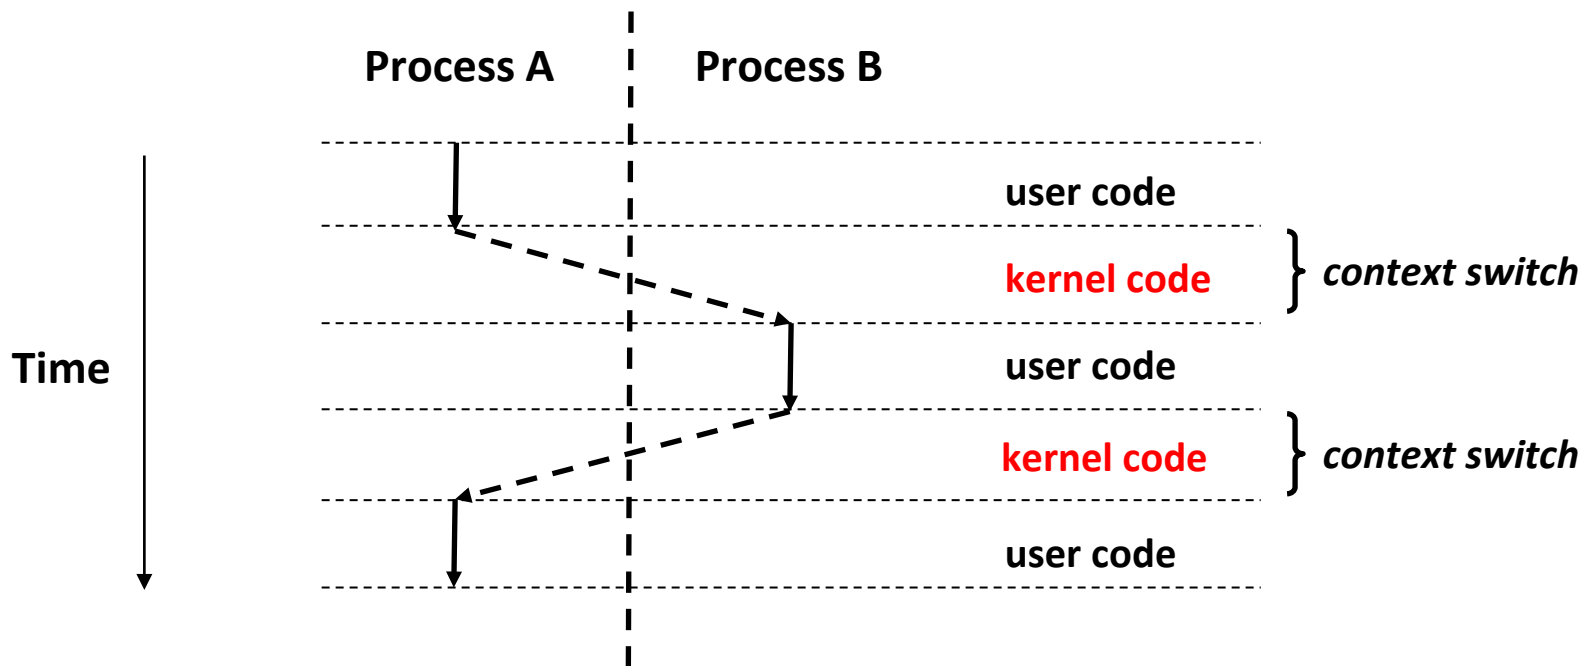
Logical Control Flows (2)

- User view of concurrent processes
 - Control flows for concurrent processes are physically disjoint in time
 - However, we can think of concurrent processes are running in parallel with each other



Context Switching

- Control flow passes from one process to another via a **context switch**



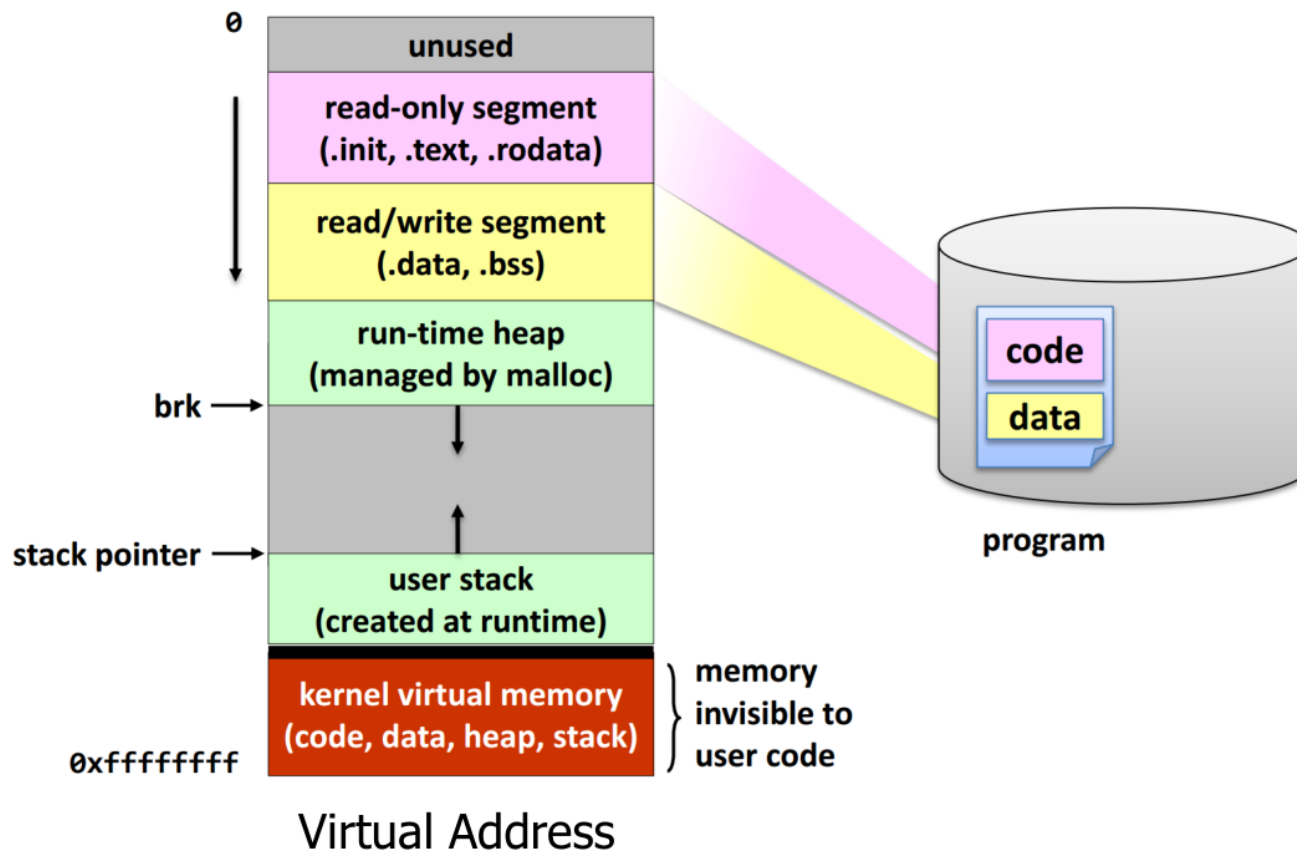
Process Control Block

- A data structure used by OS to store all the information for efficient process management

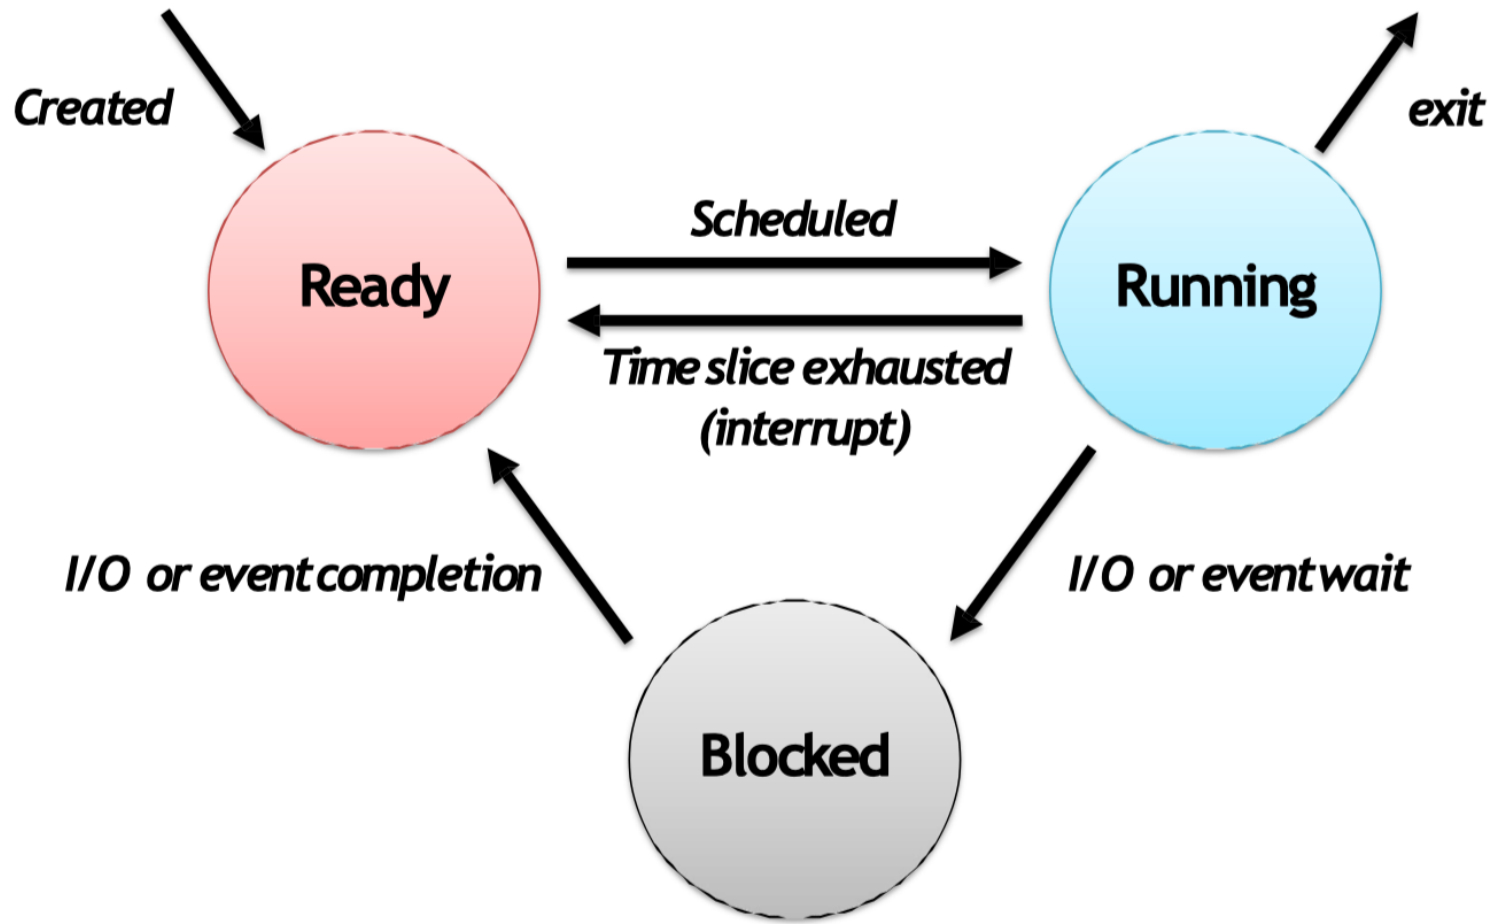
Process ID
Process State
Program Counter
Registers
Memory Info.
List of Open Files
Process Priority
I/O Status, etc...

Private Address Space

- Process in memory




Process State Transition



Process Hierarchy

- Parent-child relationship
 - One process can create another process
 - Unix calls the hierarchy a "process group"
 - Windows has no concept of process hierarchy
- Browsing a list of processes:
 - *ps* in Unix
 - Task manager (*taskmgr*) in Windows

```
$ cat file1 | wc
```



Creating a New Process

▪ `pid_t fork (void)`

- Create a new process (child process) that is identical to the calling process (parent process)
- Return 0 to the child process
- Return child's **pid** to the parent process

```
if (fork() == 0) {  
    printf("hello from child\n");  
} else {  
    printf("hello from parent\n");  
}
```

Fork is interesting
(and often confusing)
because it is called
once but returns *twice*

Fork Example (1)


- Key points
 - Parent and child both run **same code**
 - Distinguish parent from child by **return value** from **fork()**
 - Start with same state, but each has private copy
 - Share file descriptors, since child inherits all open files

```
void fork1()
{
    int x = 1;
    pid_t pid = fork();
    if (pid == 0) {
        printf("Child has x = %d\n", ++x);
    } else {
        printf("Parent has x = %d\n", --x);
    }
    printf("Bye from process %d with x = %d\n", getpid(), x);
}
```

Fork Example (2)

- Key points
 - Both parent and child can continue forking

```
void fork2()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```



Destroying a Process

- **void exit (int status)**
 - Exit a process
 - Normally returns with status 0
 - **atexit()** registers functions to be executed upon exit

```
void cleanup(void)
{
    printf("cleaning up\n");
}

void fork3()
{
    atexit(cleanup);
    fork();
    exit(0);
}
```


Synchronizing with Children

- **pid_t wait (int *status)**
 - Suspend the current process until one of its children terminates
 - Return value is the **pid** of the child process terminated
 - if **status != NULL**, then the object it points to will be set to a status indicating why the child process terminated
- **pid_t waitpid (pid_t pid, int *status, int options)**
 - Can wait for specific process
 - Various options

Wait Example (1)

```
void fork4()
{
    int child_status;

    if (fork() == 0) {
        printf("HC: hello from child\n");
    } else {
        printf("HP: hello from parent\n");
        wait(&child_status);
        printf("CT: child has terminated\n");
    }
    printf("Bye\n");
    exit();
}
```



Wait Example (2)

- If multiple children completed,
 - Will take in arbitrary order
 - Can use macros **WIFEXITED()** and **WEXITSTATUS()** to get information about exit status

```
void fork5()
{
    pid_t pid[N];
    int i, child_status;
    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* Child */
    for (i = 0; i < N; i++) {
        pid_t wpid = wait(&child_status);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminate abnormally\n", wpid);
    }
}
```

Waitpid Example

```
void fork6()
{
    pid_t pid[N];
    int i, child_status;
    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* Child */
    for (i = 0; i < N; i++) {
        pid_t wpid = waitpid(pid[i], &child_status, 0);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminated abnormally\n", wpid);
    }
}
```

Process Termination

- Normal / error exit (voluntary)
- Fatal error (involuntary)
 - Segmentation fault (illegal memory access)
 - Protection fault
 - Exceed allocated resources, etc.
- Killed by another process (involuntary)
 - By receiving a signal
- **Zombie** process: terminated, but not removed

Zombie (1)

- Zombie process
 - Living corpse, half alive and half dead
 - When a process terminates, still consumes system resources
 - Various tables maintained by OS
- Reaping
 - Performed by parent on terminated child
 - Parent is given exit status information
 - Kernel discards the terminated process
- What if parent doesn't reap?
 - If any parent terminates without reaping a child, then child will be reaped by `init` process
 - Only need explicit reaping for long-running processes
 - e.g. shells and servers

Zombie (2)

```
linux> ./fork7 &
[1] 6639
Running Parent, PID = 6639
Terminating Child, PID = 6640
linux> ps
  PID TTY          TIME CMD
 6585 ttyp9        00:00:00 tcsh
 6639 ttyp9        00:00:03 fork7
 6640 ttyp9        00:00:00 fork7 <defunct>
 6641 ttyp9        00:00:00 ps
linux> kill 6639
[1] Terminated
linux> ps
  PID TTY          TIME CMD
 6585 ttyp9        00:00:00 tcsh
 6642 ttyp9        00:00:00 ps
```

```
void fork7()
{
    if (fork() == 0) {
        /* Child */
        printf("Terminating Child, PID = %d\n",
               getpid());
        exit(0);
    } else {
        printf("Running Parent, PID = %d\n",
               getpid());
        while (1); /* Infinite loop */
    }
}
```

- **ps** shows child processes as "defunct"
- Killing parent allows child to be reaped

Running New Programs (1)

- **int execl (char *path, char *arg0, ..., NULL)**
 - Load and run executable at **path** with arguments **arg0, arg1, ...**
 - **path** is the complete path of an executable
 - **arg0** becomes the name of the process
 - Typically **arg0** is either identical to **path**, or else it contains only the executable filename from path
 - "Real" arguments to the executable start with **arg1**, etc.
 - List of args is terminated by NULL
 - Return -1 if error, otherwise doesn't return!
- **int execv (char *path, char *argv[])**
 - argv: null terminated pointer arrays

Running New Programs (2)

- Example: running /bin/ls

```
main()
{
    if (fork() == 0) {
        execl("/bin/ls", "ls", "/", NULL);
    }
    wait(NULL);
    printf("completed\n");
    exit();
}
```

```
main()
{
    char *args[] = {"ls", "/", NULL};
    if (fork() == 0) {
        execv("/bin/ls", args);
    }
    wait(NULL);
}
```

Summary

- Process abstraction
 - Logical control flow
 - Private address space
- Process-related system calls
 - fork()
 - exit(), atexit()
 - wait(), waitpid()
 - execl(), execle(), execv(), execve(), ...

Exercise

- Make simple mini shell
 - This program should be under an infinite loop with **conditional exit** ("quit")
 - When the command is entered, the command is executed using the **child process**
 - When the program is quitted, the parent process must wait for all the child processes to be done before exiting itself
 - The mini shell only executes programs under **/bin** directory
 - There is no limit to the header file

```
root@swin:/home/swe2024# ./ex3
ls
ex3 ex3.c test
echo test
test
cat test
example3
```

Skeleton code of p5.c

- copy the skeleton code to your directory

```
$ cp ~swe2024-41_23s/2023s/p5_skeleton.c/
```

```
#!/bin/bash

# 1) compile your program with makefile created in Exercise 1
# your code here

# 2) make answer sheet as answer.txt
# your code here

# 3) run your program and save output as output.txt
echo "22 11" | ./p2.out > output.txt

# 4) compare answer.txt and output.txt with diff command as result.txt
# your code here

# 5) print the contents of "result.txt"
# your code here

# 6) clean object files and executable file
# your code here
```

Exercise

- Submit your Makefile & p5.c as p5
 - InUiYeJi Cluster
 - Submit the folder into p5
 - `~swe2024-41_23s/bin/submit p5 p5`
 - Due

```
2014311240@swji:/home/2014311240$ ~swe2024-41_23s/bin/submit p5 p5
user name           :2014311240
Submitted Files for p5:
File Name                                File Size      Time
-----
p5-2014311240-Mar.29.17.12.153059682    730            Wed Mar 29 17:12:53 2023
2014311240@swji:/home/2014311240$
```

./p5


- Makefile
- p5.c

Exercise Hint

- *getline* function

```
char *cmd;  
size_t size;  
getline(&cmd, &size, stdin);  
cmd[strlen(cmd)-1] = '\\0';  
free(cmd);
```

cat test



c	a	t		t	e	s	t	\\n
c	a	t		t	e	s	t	\\0

- *strtok* function

```
int i = 0;  
char *ptr = strtok(cmd, " ");  
  
while (ptr != NULL) {  
    arg[i++] = ptr;  
    ptr = strtok(NULL, " ");  
}  
arg[i] = NULL;
```

Exercise Hint

- *sprintf()* function

```
char path[100];  
sprintf(path, "/bin/%s", arg[0]);
```

/	b	i	n	/	c	a	t	\0
---	---	---	---	---	---	---	---	----

– Alternative: *snprintf()* for safety

Summary Report

- Summary report about man command result of
 - `fork()`
 - `exit()`
 - `wait()`
 - `execv()`
- Submission form
 - A4 size PDF format (No page limitation)
 - [SWE2024 Report-5] studentID_name
 - Ex) [SWE2024 Report-5] 2022XXXXXX_박재형
 - Submit to iCampus
 - Due: until Friday, 31 March 2023, 23:59