

Pipe and Redirection

Prof. Joonwon Lee (joonwon@skku.edu)

TA – Jaehyung Park (jaeseanpark@gmail.com)

TA – Luke Albano (lukealbano@arcs.skku.edu)

Sungkyunkwan University

Everything is a File

- Actually, “Everything is a file descriptor”
- Pros
 - Can reuse tools, APIs on a wide range of resources
- Cons
 - Not a fast approach
- Communication using file interface?

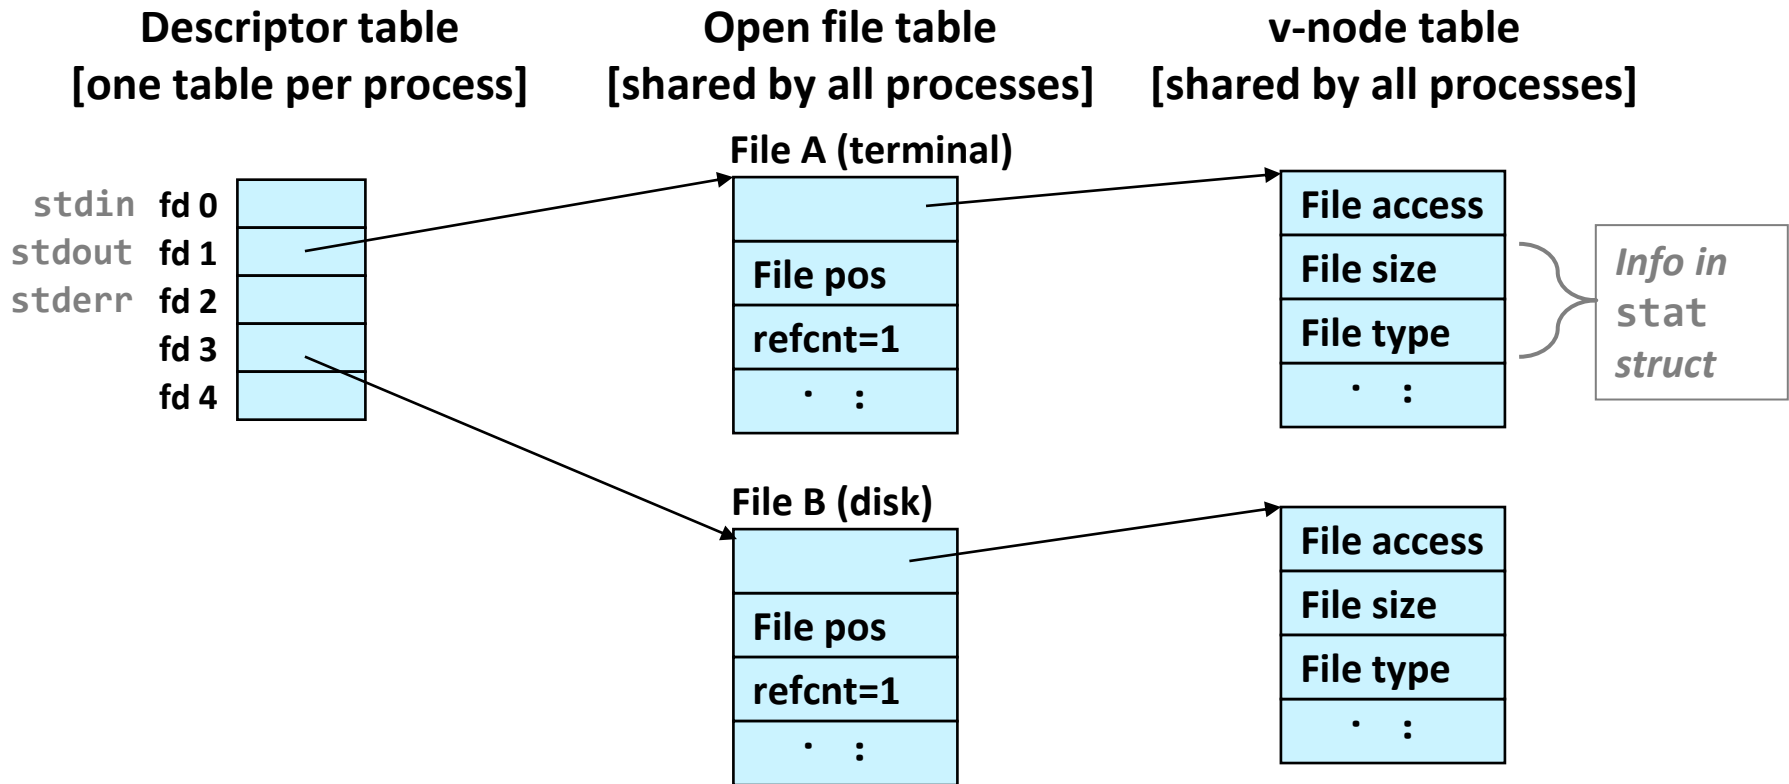
- Inter-process Communication
 - Mechanism for communication between processes
- Methods in IPC
 - Signals
 - I/O redirection
 - Anonymous Pipe
 - Named Pipe (FIFO)
 - Shared Memory, Message Queue, Message Passing, etc.

Open Files in Kernel (1)

- How the Unix kernel represents open files?
- 3-levels
 - Descriptor table
 - 1 table per process
 - Pointer to entry in the “file table”
 - Open file table
 - Shared by all processes
 - Current file position, reference count, pointer to entry in the “v-node table”
 - v-node table
 - Shared by all processes
 - Information about file itself (size, permission, ...)

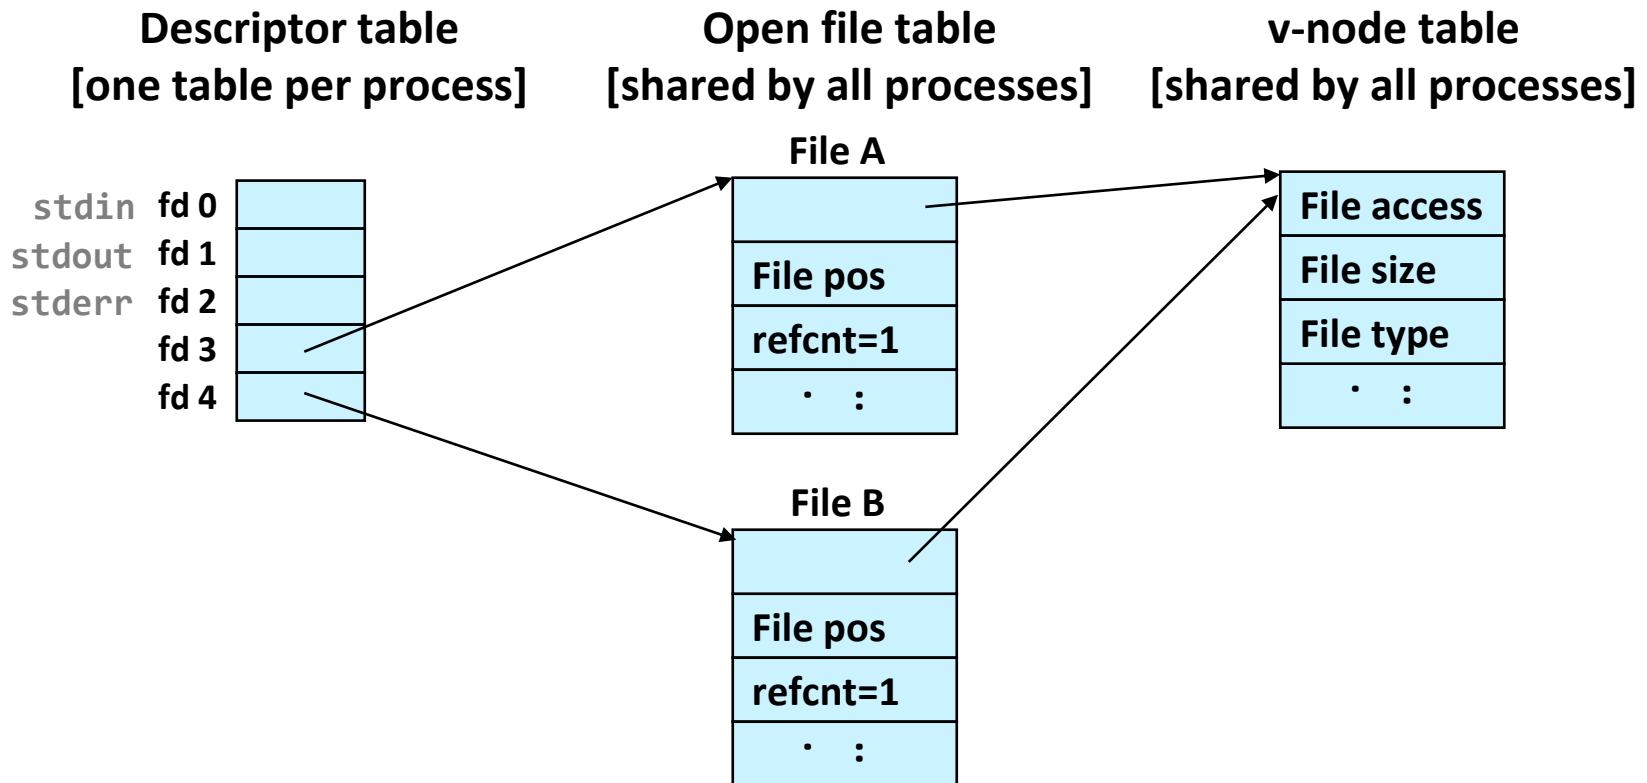
Open Files in Kernel (2)

- How the Unix kernel represents open files?



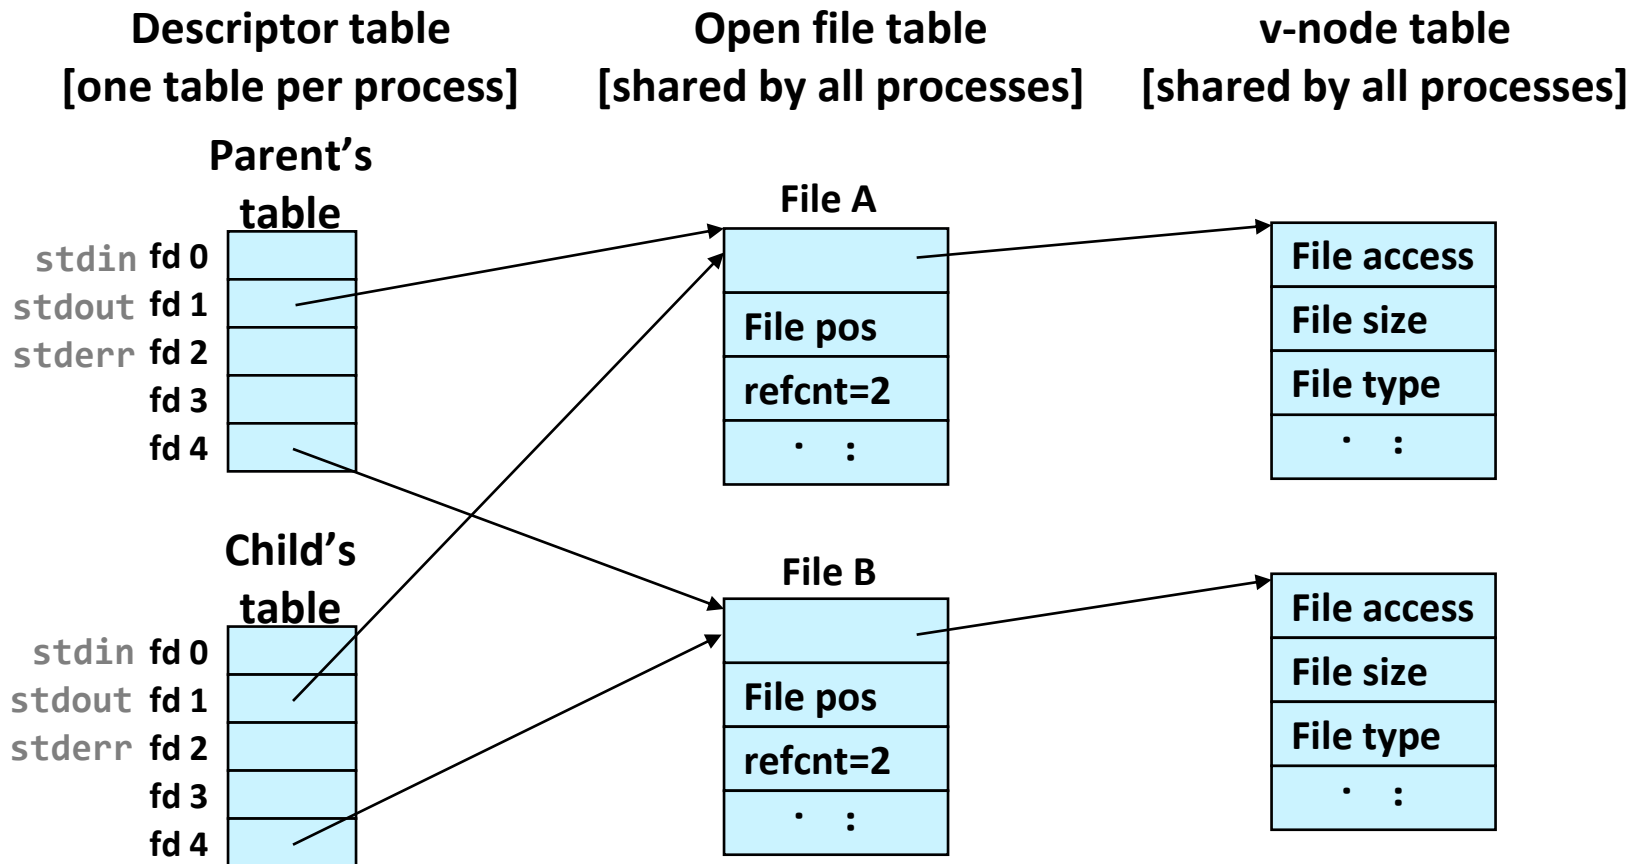
Open Files in Kernel (3)

- Calling `open()` twice with the same filename



Open Files in Kernel (4)

- Calling `fork()`



Open Files in Kernel (5)

- What will be the result?

```
// ~swe2024-41_23s/2023s/w8/open_file.c

#include <unistd.h>
#include <fcntl.h>
#include <stdlib.h>
#include <sys/wait.h>
#include <assert.h>
#include <stdio.h>

int main(void) {
    char c;
    int fd = open("temp.txt", O_RDONLY);

    if (fork() == 0) { // child process
        read(fd, &c, 1); // reads 1 character
        exit(0); // child terminates
    } else {
        wait(NULL); // waits for child to exit
        read(fd, &c, 1); // reads 1 character
        printf("c=%c\n", c); // prints read character
    }

    return 0;
}
```

temp.txt


sysexperiment

I/O Redirection

- Q: How does a shell implement I/O redirection?
 - \$ ls > foo.txt
- A: By calling the dup2(oldfd, newfd) function.
 - Copies (per-process) descriptor table entry **oldfd** to entry **newfd**

Descriptor table
before dup2(4,1)

fd 0	
fd 1	a
fd 2	
fd 3	
fd 4	b

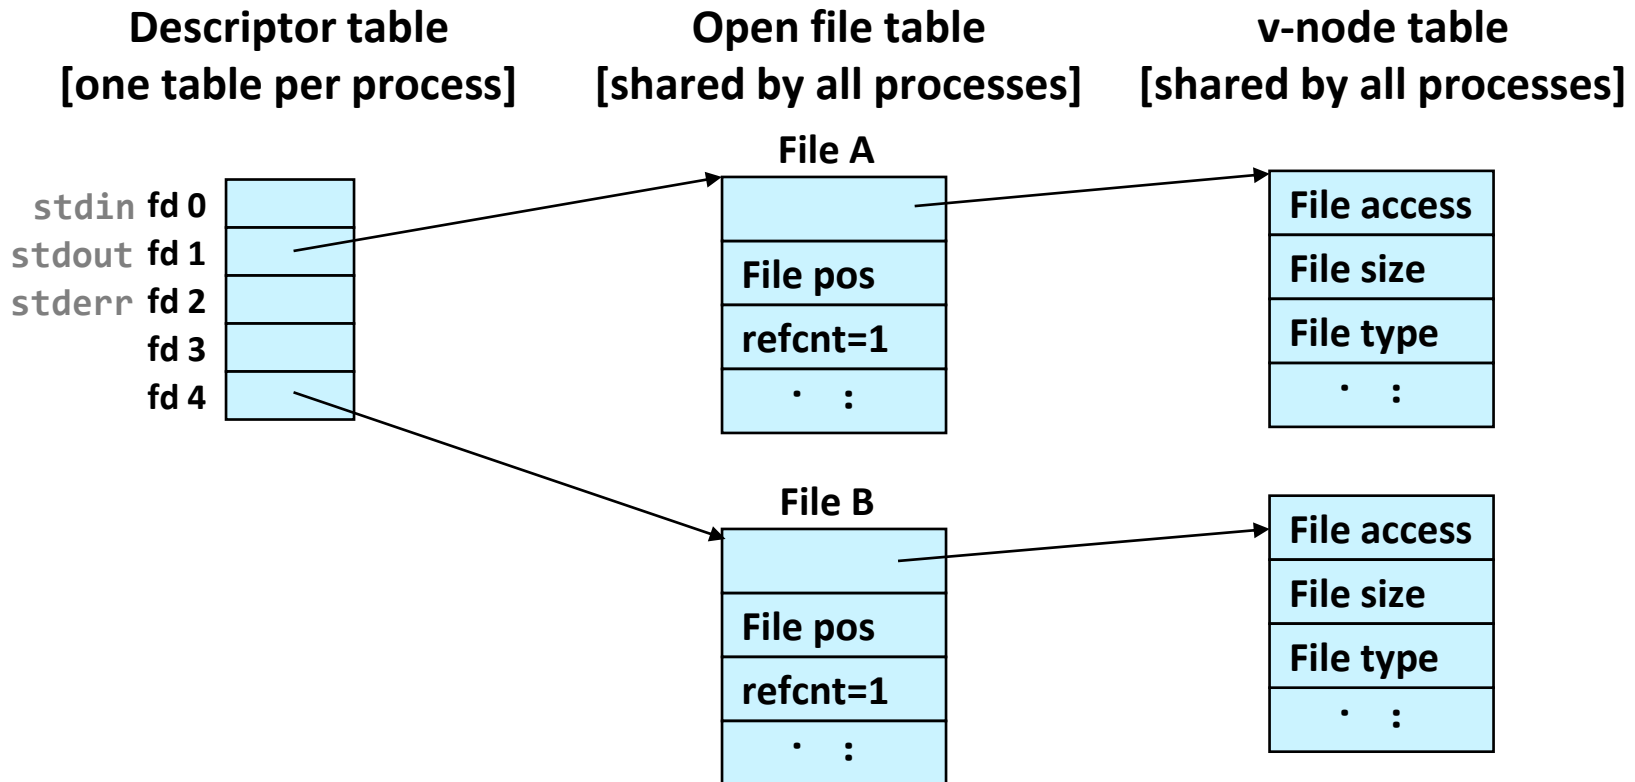


Descriptor table
after dup2(4,1)

fd 0	
fd 1	b
fd 2	
fd 3	
fd 4	b

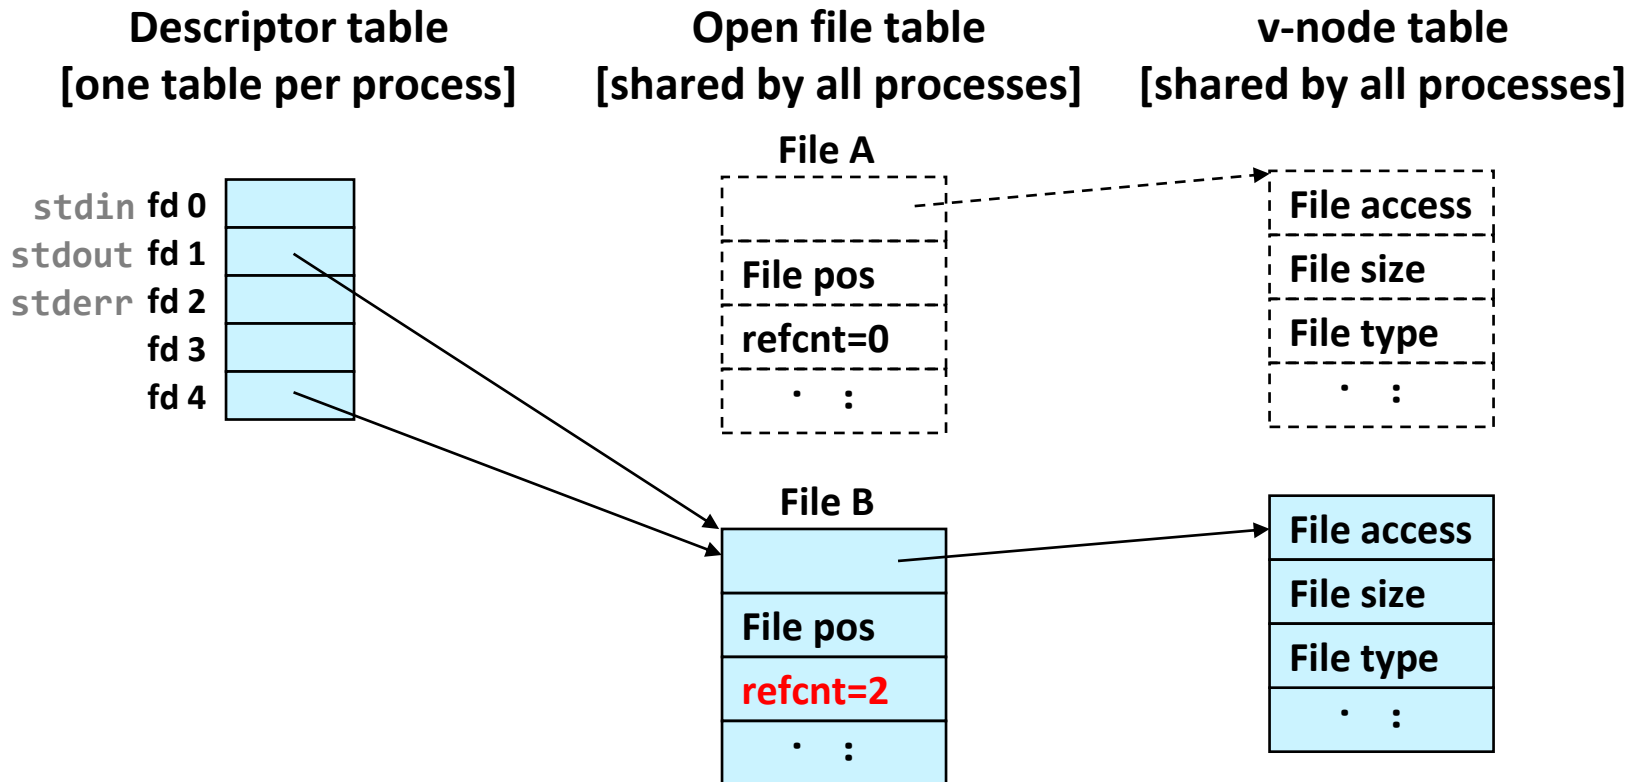
I/O Redirection Example (1)

- Before calling `dup2(4,1)`



I/O Redirection Example (2)

- After calling `dup2(4,1)`




Pipes

- Pipes
 - Oldest mechanism for IPC in UNIX
 - Communicate with shared file descriptor of pipe
- Limitations
 - Half-duplex: data flows only in one direction
 - Data only can be read once
- Two pipes
 - Anonymous pipe
 - Temporary, between parent-child
 - Named pipe
 - Saved in file system, between arbitrary processes


Anonymous Pipe (1)

- `int pipe (int fd[2])`
 - Two file descriptors are returned through the **fd** argument
 - **fd[0]**: open for reading
 - **fd[1]**: open for writing
 - The output of **fd[1]** is the input for **fd[0]**
 - On error, return -1




Anonymous Pipe (2)

**parent => child:
parent closes fd[0];
child closes fd[1];**



Anonymous Pipe (3)

**parent \leq child:
parent closes fd[1];
child closes fd[0];**



Using Anonymous Pipe

```
// ~swe2024-41_23s/2023s/w8/anonymous_pipe.c
#include <unistd.h>
#include <stdlib.h>
#define MAXLINE 80

int main(void) {
    int n, fd[2]; // fd is used to receive the file descriptor of pipe
    pid_t pid;
    char line[MAXLINE]; // Contains a "string" of MAXLINE characters
    if(pipe(fd) < 0) exit(1); // Create an anonymos pipe, exit if failed
    if((pid = fork()) < 0) exit(2); // Fork to create child process, exit if
    failed

    if (pid > 0) { // parent
        close(fd[0]); // Close as this fd will not be used by parent
        write(fd[1], "hello world\n", 12); // Writes "hello world" to fd[1]
    } else { // child
        close(fd[1]); // Close as this fd will not be used by child
        n = read(fd[0], line, MAXLINE); // Receives "hello world" from fd[0]
        write(1, line, n); // Writes "hello world" to fd 1 (STDOUT)
    }
    exit(0);
}
```


Named Pipe (FIFO)


- `int mknod (const char *path, mode_t mode, dev_t dev)`
- Ex) `mknod ("path", S_IFIFO | 0666, 0);`
- `int mkfifo (const char *path, mode_t mode)`
- Ex) `mkfifo ("path", 0660);`
- You can also make FIFOs on the command line by using "mknod" and "mkfifo"

FIFO Example (1)

- Duplicating a stream

- Passing data from one shell pipeline to another without creating intermediate temporary files

```
$ mkfifo fifo1  
$ cat < fifo1 > output.txt &  
$ cat < input.txt > fifo1
```




FIFO Example (1)

- Duplicating a stream

- Passing data from one shell pipeline to another without creating intermediate temporary files

```
$ mkfifo fifo1  
$ prog3 < fifo1 &  
$ prog1 < infile | tee fifo1 | prog2
```




FIFO Example (2)

■ Client-server communication

– A client-server application to pass data between the client and server **on the same machine**

- Clients write to a “well-known” FIFO to send a request to the server



Using FIFOs

- Opening a FIFO
 - An open for read (write)-only blocks until some other process opens the FIFO for writing (reading)
- Reading / Writing a FIFO
 - Writing to a FIFO that no process has open for reading causes SIGPIPE to generate
 - When the last writer for a FIFO closes the FIFO, an end of file is generated for the reader of the FIFO
 - PIPE_BUF: the maximum amount of data that can be written atomically to a FIFO (without being interleaved among multiple writers)

Lab Exercise

- Make *mini shell* which supports **I/O redirection & pipes**
 - This program should be under an infinite loop with **conditional exit** (“quit”)
 - When the command is entered, the command is executed using the **child process**
 - When the program is quitted, the parent process must wait for all the child processes to be done before exiting itself
 - The *mini shell* only executes programs under **/bin** directory
 - The *mini shell* must handle **2 types of redirections(>, <)** and **pipe(|)**
 - There is no limit to the header file

Lab Exercise

- Skeleton code
 - `cp ~/swe2024-41_23s/2023s/p8_skeleton.c ./`
- There will be **only one** IPC at one line
 - (e.g. No command such as `: cat < test.txt | grep test`)

- Example)

```
sw@SW:~/swe2024/week8/exercise$ ./ex8
$ ls
ex8  ex8.c
$ echo test > t.txt
$ cat t.txt
test
$ grep include < ex8.c
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <stdlib.h>
#include <string.h>
#include <sys/types.h>
#include <sys/wait.h>
$ ls -la | grep ex
$ -rwxrwxr-x 1 sw sw 17400 10월 22 18:14 ex8
-rw-rw-rw- 1 sw sw 3316 10월 22 18:14 ex8.c
ls
ex8  ex8.c  t.txt
$ quit
sw@SW:~/swe2024/week8/exercise$
```

Exercise Submission

- Submit your exercise code and Makefile
 - InUiYeJi cluster
 - Submit the folder into p8
 - `~swe2024-41_23s/bin/submit p8 p8`
 - Due date: Sunday, 23 April 2023, 23:59
 - We will compile by using command *make*
 - If compilation fails, your points for this exercise will be **zero**

Summary Report

- Summary report about man command result of
 - pipe()
 - mkfifo()
- Submission form
 - A4 size PDF format (No page limitation)
 - [SWE2024 Report-8] studentID_name
e.g.) [SWE2024 Report-8] 2023XXXXXX_홍길동
 - Submit to iCampus
 - Due: until Sunday, 23 April 2023, 23:59