

# System-V IPC

---

Prof. Joonwon Lee ([joonwon@skku.edu](mailto:joonwon@skku.edu))

TA – Jaehyung Park ([jaeseanpark@gmail.com](mailto:jaeseanpark@gmail.com))

TA – Luke Albano ([lukealbano@arcs.skku.edu](mailto:lukealbano@arcs.skku.edu))

Sungkyunkwan University

# Review of Inter-Process Communication

- Message passing model

- Pipe
  - Named pipe (FIFO)
  - Message queue
- Week 8 (Pipe & Redirection)

- Shared memory model

- Shared memory

# System V IPC Interface

- System V IPC provides
  - Shared memory, message queue, semaphores, ...
- Key & Identifier
  - Key: identify user (process) **like file name**
    - Key can be hard coded or obtained by `ftok()` function
  - Identifier: uniquely identify an **IPC object**
- Permission
  - XYZ form
  - Like file permission
  - X: owner user, Y: owner group, Z: every other
  - $100_2$ : read (-r)       $010_2$ : write (-w)       $001_2$ : modify (-a)

# System V IPC API

## ■ `key_t ftok (const char* pathname, int proj_id)`

- To create deterministic key with string
- **pathname**: any valid file path
- **proj\_id**: some integer must be non-zero  
(only lower 8 bits are used)

```
int main(void)
{
    key_t k1 = ftok(".", 'a');
    printf("key1:\t%d\n", k1);
    key_t k2 = ftok(".", 'a' + 1);
    printf("key2:\t%d\n", k2);

    return 0;
}
```

```
key1: 1627521876
key2: 1644299092
```

**NAME**  
ftok - convert a pathname and a project identifier to a System V IPC key

**SYNOPSIS**  
`#include <sys/types.h>`  
`#include <sys/ipc.h>`  
  
`key_t ftok(const char *pathname, int proj_id);`

**DESCRIPTION**  
The `ftok()` function uses the identity of the file named by the given `pathname` (which must refer to an existing, accessible file) and the least significant 8 bits of `proj_id` (which must be nonzero) to generate a `key_t` type System V IPC key, suitable for use with `msgget(2)`, `semget(2)`, or `shmget(2)`.  
  
The resulting value is the same for all pathnames that name the same file, when the same value of `proj_id` is used. The value returned should be different when the (simultaneously existing) files or the project IDs differ.

**RETURN VALUE**  
On success, the generated `key_t` value is returned. On failure -1 is returned, with `errno` indicating the error as for the `stat(2)` system call.

# System V IPC Interface Data Structure

## ■ ipc\_perm

- Data structure designed to handle permission
- Component: user and group ID, mode bit, and sequence number

```
/* defined <sys/ipc.h> */

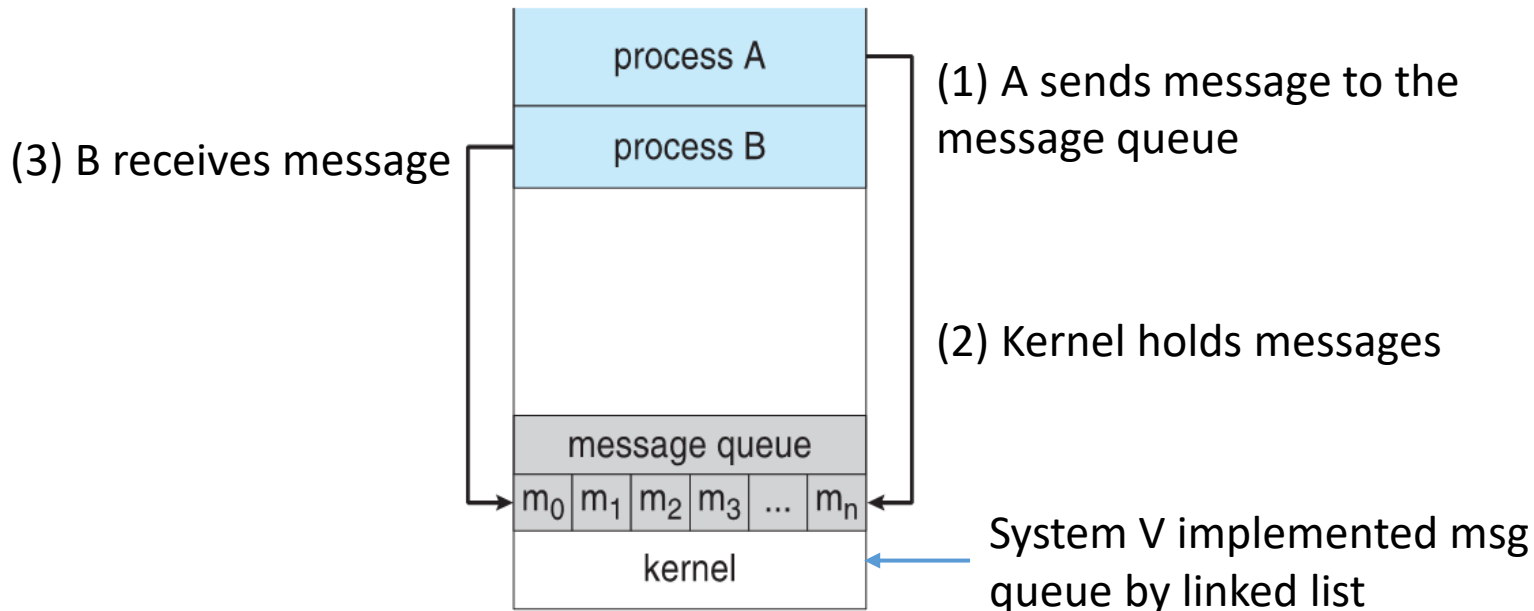
struct ipc_perm {
    key_t      key;      /* user specified msg/sem/shm key */
    uid_t      cuid;     /* creator user id */
    gid_t      cgid;     /* creator group id */
    uid_t      uid;      /* user id */
    gid_t      gid;      /* group id */
    mode_t      mode;     /* r/w permission */
    unsigned short seq;   /* sequence # */
};
```

# Message Queue

- Mail-box like communication method

- Pros: easy to use
- Cons: overhead if data is large

- Typical message passing model



# Message Queue Data Structure

- Defined in `<sys/msg.h>`

	<code>struct msqid_ds {</code>	
	<code>struct ipc_perm msg_perm;</code>	<code>/* permission */</code>
as a list	<code>struct msg *msg_first;</code>	<code>/* first message */</code>
	<code>struct msg *msg_last;</code>	<code>/* last message */</code>
resource	<code>unsigned long msg_cbytes;</code>	<code>/* number of bytes in use */</code>
	<code>unsigned long msg_qnum;</code>	<code>/* number of msgs */</code>
	<code>unsigned long msg_qbytes;</code>	<code>/* max # of bytes */</code>
access	<code>pid_t msg_lspid;</code>	<code>/* pid of last sender() */</code>
	<code>pid_t msg_lrpid;</code>	<code>/* pid of last receiver() */</code>
	<code>time_t msg_stime;</code>	<code>/* time of last send() */</code>
	<code>time_t msg_rtime;</code>	<code>/* time of last receive() */</code>
	<code>time_t msg_ctime;</code>	<code>/* time of last msgctl() */</code>
	<code>};</code>	

# Message Structure

- **msgbuf** should be implemented
  - The first attribute of **msgbuf** is **msgtype**
  - Sender and receiver should share the same **msgbuf** implementation

```
struct msgbuf{  
    long msgtype  
    ...  
};
```

Must have attribute **msgtype**  
And **msgtype** must be > 0

```
struct msgbuf_text{  
    long msgtype  
    char text[256];  
};
```

```
struct msgbuf_nums{  
    long msgtype  
    int numbers[8];  
};
```



# Message Queue API

- **int msgget (key\_t key, int msgflg)**
  - e.g., msgget (0x1234, IPC\_CREAT)
- **int msgsnd (int msgid, const void \*msgp, size\_t msgsz, int msgflg)**
  - e.g., msgsnd (0x1234, (void \*)buf, sizeof(struct msgbuf), 0)
- **ssize\_t msgrcv (int msgid, void \*msgp, size\_t msgsz, long msgtyp, int msgflg)**
  - e.g., msgrcv (0x1234, (void \*)buf, sizeof(struct msgbuf), 1, 0)
- **int msgctl (int msgid, int cmd, struct msqid\_ds \*buf)**
  - e.g., msgctl (0x1234, IPC\_RMID, NULL)      /\* remove msgq \*/

# System V IPC Flags

- **IPC\_CREAT**
  - Like **O\_CREAT**
  - Create object if there is no object created by given key
- **IPC\_EXCL**
  - Like **O\_EXCL**
  - Return error if object is already exists
- **IPC\_SET**
  - Set data of the object with given ID
  - Parameter of **ctl()** function
    - e.g., **msgctl(0x1234, IPC\_SET, ptr)**

# System V IPC Flags

## ■ IPC\_RMID

- To remove object with given ID
- Parameter of **ctl()** function
  - e.g., `msgctl(0x1234, IPC_RMID, NULL)`

## ■ IPC\_STAT

- Get data of the object with given ID
- Return object specific data structure
  - e.g., `msgqid_ds`, `shmid_ds`
- Parameter of **ctl()** function
  - e.g., `msgctl(0x1234, IPC_STAT, ptr)`

## ■ IPC\_NOWAIT

- Return immediately if no message of the requested type is in the queue. The system call fails with `errno` set to `ENOMSG`.

# Message Queue Sender Example

```
int main(void)
{
    key_t k = ftok(".", 'a');           // get key
    int qid = msgget(k, IPC_CREAT | 0666); // get msg queue object
    if (qid == -1) {
        perror("msgget error : ");
        exit(1);
    }

    struct msgbuf *msg = malloc(sizeof(struct msgbuf));
    msg->msgtype = 1;           // set msgtype 1
    strcpy(msg->text, "Hello world?\n"); // write message

    /* send */
    if (msgsnd(qid, (void *)msg, sizeof(struct msgbuf), 0) == -1) {
        perror("msgsnd error : ");
        exit(1);
    }

    free(msg);
    return 0;
}
```

```
struct msgbuf {
    long msgtype;
    char text[256];
};
```

# Message Queue Receiver Example

```
int main(void)
{
    key_t k = ftok(".", 'a');           // get key
    int qid = msgget(k, IPC_EXCL | 0666); // get msg queue object
    if (qid == -1) {
        perror("msgget error : ");
        exit(1);
    }

    struct msgbuf *msg = malloc(sizeof(struct msgbuf));

    /* receive */
    if (msgrcv(qid, (void *)msg,
               sizeof(struct msgbuf), 1, 0) == -1) {
        perror("msgrsv error : ");
        exit(1);
    }


    printf("%s", msg->text);
    free(msg);
    return 0;
}
```

```
struct msgbuf {
    long msgtype;
    char text[256];
};
```

# Shared Memory Model

- Mapping segment of memory that will be shared by more than one process
- Fastest form of IPC
  - There is no intermediation
- Shared memory can be written to and read from by any number of process
  - Pros: fast
  - Cons: synchronization

# Shared Memory Abstract



# Shared Memory API

- **int shmget (key\_t key, size\_t size, int shmflag)**
  - e.g., **shmget (0x1234, sizeof(int) \* 100, IPC\_CREAT)**
  - Use flags start with SHM\_HUGE to allocate large size memory
- **void \*shmat (int shmid,**  
**const void \*shmaddr, int shmflg)**
  - e.g., **shmat(1, NULL, 0)**
  - **shmaddr**: usually NULL
  - **shmflag**: set 0 for READ WRITE permission



# Shared Memory API (Cont.)

- **int shmdt (int shmid,**  
**const void \* shmaddr, int shmflag)**
  - e.g., **shmdt (1, shmaddr, 0)**
  - **shmaddr**: address where shared memory segment is attached
- **void shmctl (int shmid,**  
**int cmd, struct shmid\_ds \*buf)**
  - e.g., **shmctl(1, IPC\_RMID, 0)** */\* remove ID 1 \*/*
  - See common IPC flags

# Shared Memory Data Structure

- Defined in <sys/shm.h>

```
struct shmid_ds {
    struct ipc_perm    shm_perm;    /* operation perms */
    int                shm_segsz;    /* size of seg (bytes) */
    time_t             shm_atime;    /* last attach time */
    time_t             shm_dtime;    /* last detach time */
    time_t             shm_ctime;    /* last change time */
    pid_t              shm_cpid;     /* pid of creator */
    pid_t              shm_lpid;     /* pid of last operator */
    unsigned short     shm_nattch;   /* no. of curr attaches */

    /* ----- unused ----- */
    unsigned short     shm_unused;
    void               *shm_unused2;
    void               *shm_unused3;
};
```

# Shared Memory Example

```
int main(void)
{
    key_t k = ftok(".", 'b');
    int shm_id = shmget(k, sizeof(int), IPC_CREAT | 0666);
    if (shm_id < 0) {
        perror("shmget fail");
        exit(0);
    }

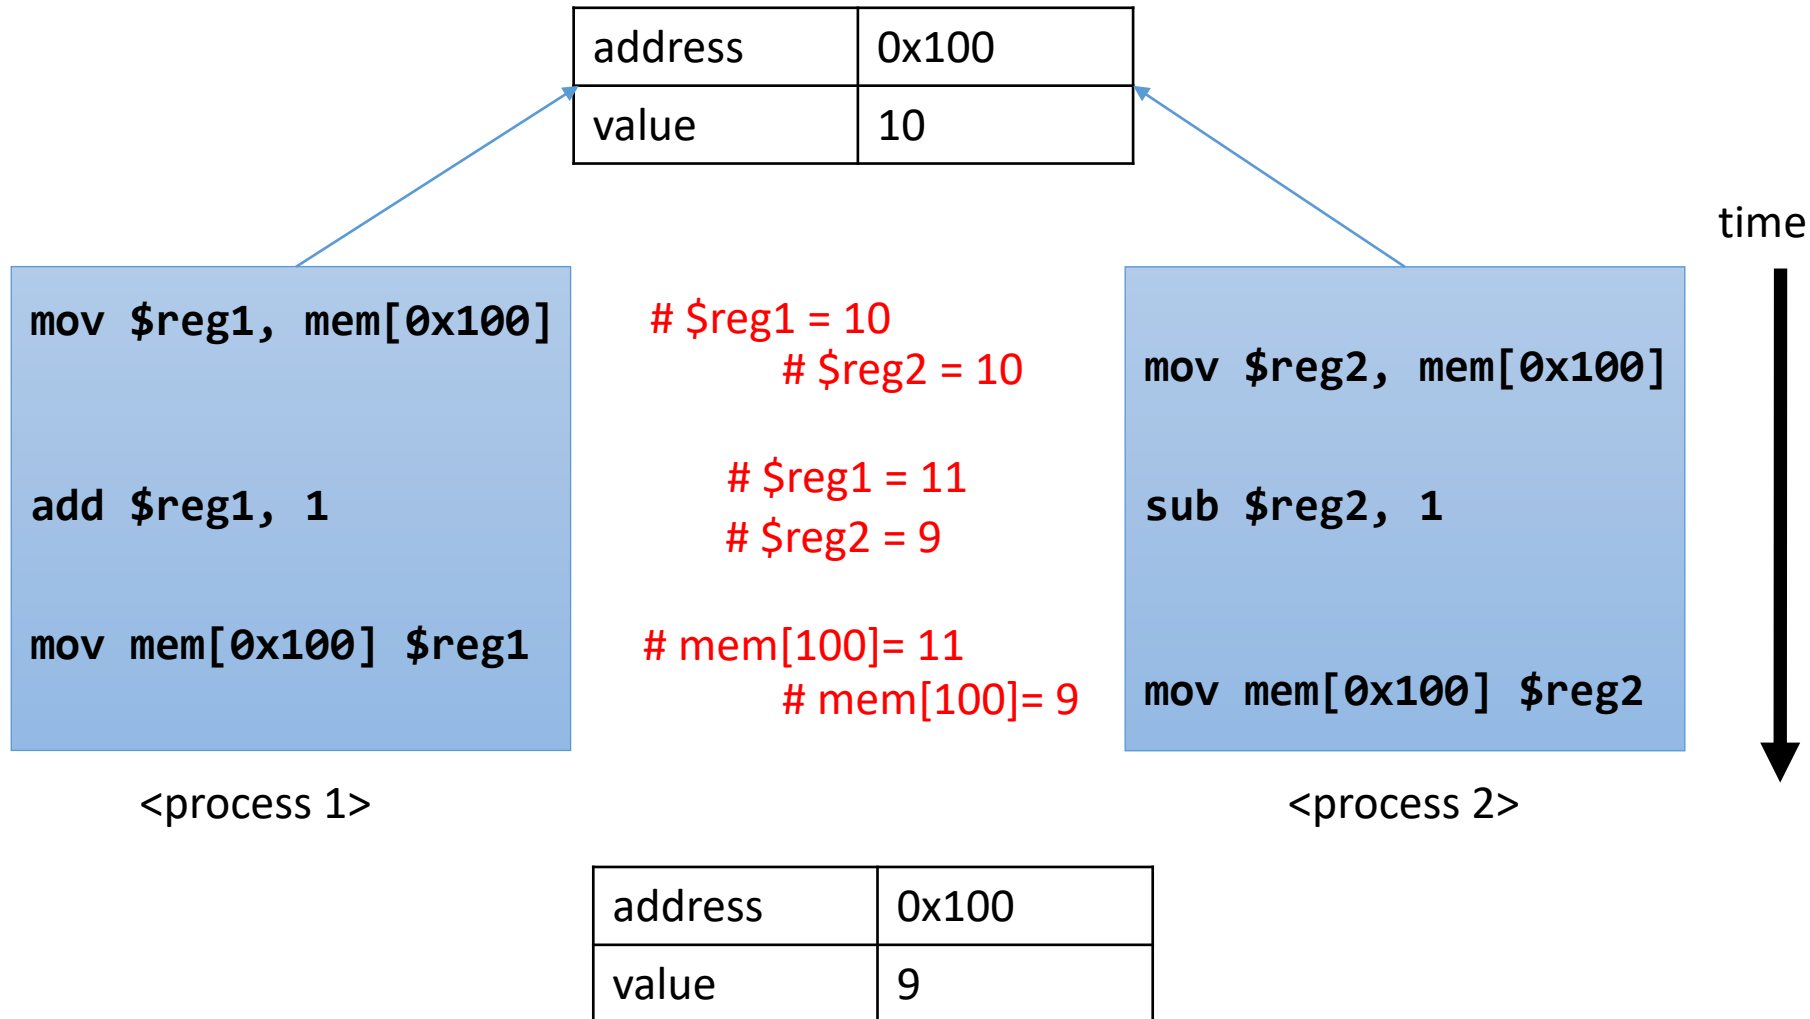
    int *shmaddr = shmat(shm_id, NULL, 0);
    *shmaddr = 1;
    printf("value before fork: %d\n", *shmaddr);

    if (fork() == 0) {
        (*shmaddr) += 1;
        printf("value in the child: %d\n", *shmaddr);
        exit(0);
    } else {
        wait(NULL);
        printf("value in the fork: %d\n", *shmaddr);
    }

    return 0;
}
```

```
jaehyun@cs1:~/SSE/week9$ ./shm
value before fork: 1
value in the child: 2
value in the fork: 2
```

# Race Condition



# Race Condition Example

```
int main(void)
{
    key_t k = ftok(".", 'b');
    int shm_id = shmget(k, sizeof(int), IPC_CREAT | 0666);
    if (shm_id < 0) {
        perror("shmget fail");
        exit(0);
    }

    int *shmaddr = shmat(shm_id, NULL, 0);
    *shmaddr = 1;
    printf("value before fork: %d\n", *shmaddr);

    if (fork() == 0) {
        for (int i = 0; i < 1000000; i++) /* Child add 1M */
            (*shmaddr) += 1;
        exit(0);
    } else {
        for (int i = 0; i < 1000000; i++) /* Parent sub 1M */
            (*shmaddr) -= 1;
    }
    printf("value after operation: %d\n", *shmaddr); /* Must 1 */
    return 0;
}
```

```
jaehyun@cs1:~/SSE/week9$ ./race
value before fork: 1
value after operation: 519688
jaehyun@cs1:~/SSE/week9$ ./race
value before fork: 1
value after operation: -813183
jaehyun@cs1:~/SSE/week9$ ./race
value before fork: 1
value after operation: -999818
jaehyun@cs1:~/SSE/week9$ ./race
value before fork: 1
value after operation: -814614
```

# Race Condition Solution Example

```
int main(void)
{
    key_t k = ftok(".", 'b');
    int shm_id = shmget(k, sizeof(int), IPC_CREAT | 0666);
    if (shm_id < 0) {
        perror("shmget fail");
        exit(0);
    }

    int *shmaddr = shmat(shm_id, NULL, 0);
    *shmaddr = 1;
    printf("value before fork: %d\n", *shmaddr);

    if (fork() == 0) {
        for (int i = 0; i < 1000000; i++) /* Child add 1M */
            (*shmaddr) += 1;
        exit(0);
    } else {
        wait(NULL);
        for (int i = 0; i < 1000000; i++) /* Parent sub 1M */
            (*shmaddr) -= 1;
    }
    printf("value after operation: %d\n", *shmaddr); /* Must 1 */
    return 0;
}
```

```
jaehyun@cs1:~/SSE/week9$ ./race.
value before fork: 1
value after operation: 1
jaehyun@cs1:~/SSE/week9$ ./race.
value before fork: 1
value after operation: 1
jaehyun@cs1:~/SSE/week9$ ./race.
value before fork: 1
value after operation: 1
jaehyun@cs1:~/SSE/week9$ ./race.
value before fork: 1
value after operation: 1
```

# Shell Commands

## ■ **ipcs (-i, -m, -q, -a)**

- List IPC objects
- **-i <ID>** : show details about given ID
- **-m** : show shared memory
- **-q** : show message queue
- **-a** : show all

```
> ipcs
----- Message Queues -----
key      msqid    owner    perms    used-bytes  messages
0x6120b818 0        jaepark  666      0            0

----- Shared Memory Segments -----
key      shmid    owner    perms    bytes       nattch     status

----- Semaphore Arrays -----
key      semid    owner    perms    nsems
```

## ■ **ipcrm (-M, -m, -Q, -q, -a)**

- Remove IPC objects
- **-m <ID>** : remove shared memory given ID (-M <key>)
- **-q <ID>** : remove message queue given ID (-Q <key>)
- **-a** : remove all IPC objects

# Exercise

- Make chatting program
  - When program starts, get ID and receiver's ID by stdin
  - User can receive message while entering message
  - When the user get the message from the other user (normal message), print out the message
  - When the receiver (the other user) check the message, send back ack message with read time
  - There should be no wait between normal message and ack message (add `IPC_NOWAIT` flags to the `msgrcv()`)
  - To finish program, enter "quit"



# Exercise

## ■ Print format

- RECEIVED **\$USER\_MSG** /\* normal message \*/
- **\$ID** read message at: **\t\$READ\_TIME** /\* ack message \*/
- Use **ctime()** function to match **\$READ\_TIME** format

```
> ./p9
My id is: 10
Enter the receiver of your message: 20
How was your exam last week?
20 read message at:      Mon Apr 24 16:07:12 2023

RECEIVED The exams were easy. Expecting straight A's.
Good for you...
20 read message at:      Mon Apr 24 16:08:15 2023

quit
```

```
> ./p9
My id is: 20
Enter the receiver of your message: 10
RECEIVED How was your exam last week?

The exams were easy. Expecting straight A's.
10 read message at:      Mon Apr 24 16:08:04 2023

RECEIVED Good for you...

quit
```

- Copy the skeleton code to your directory  
`$ cp ~swe2024-41_23s/2023s/p9_skeleton.c ./`

# Skeleton code of p9.c

- copy the skeleton code to your directory

```
$ cp ~swe2024-41_23s/2023s/p9_skeleton.c ./
```

```
#include <signal.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/msg.h>
#include <sys/types.h>
#include <sys/wait.h>
#include <time.h>
#include <unistd.h>

#define T_MSG 20

typedef struct {
    long msgtype;
    int sender;
    char text[512];
} msgbuf;

typedef struct {
    long msgtype;
    char timestamp[512];
} msgbuf_ack;

int main() {
    key_t key = ftok(".", 'a');
    int user_id, receiver_id, qid, pid;
    printf("My id is: ");
    scanf("%d", &user_id);
    getchar();
    printf("Enter the receiver of your message: ");
    scanf("%d", &receiver_id);
    getchar();
    if ((qid = msgget(key, IPC_CREAT | 0666) < 0)) {
        perror("msgget failed\n");
        exit(1);
    }

    if ((pid = fork()) == 0) {
        while (1) {
            /* child = receives messages
            1. receive normal message using buf
            1-1. print "RECEIVED [normal message]"
            1-2. send ack message using ack (current timestamp)
            --> use T_MSG + receiver_id for ack.msgtype
            2. receive ack message using read_time
            2-1. print "[receiver_id] read message at:\t[timestamp]" */
            msgbuf buf;
            msgbuf_ack read_time, ack;
            /* your code */
        }
    } else {
        while (1) {
            /* parent = sends messages
            1. get normal message from stdin
            2. if "quit", exit both parent and child (use SIGINT)
            3. send normal message */
            msgbuf buf;
            buf.msgtype = receiver_id;
        }
    }
}
```

# Exercise

- Submit your exercise source code
  - To InUiYeJi Cluster
  - Put your Makefile and \*.c files in p9 folder
  - Submit using

```
$ ~swe2024-41_23s/bin/submit p9 p9
```
  - We will compile by using command *make*
    - When compilation fails, you get zero points
    - Compiled binary name should be “p9”
- Due 2023/4/26 23:59