

# Socket

---

Prof. Joonwon Lee ([joonwon@skku.edu](mailto:joonwon@skku.edu))

TA – Jaehyung Park ([jaeseanpark@gmail.com](mailto:jaeseanpark@gmail.com))

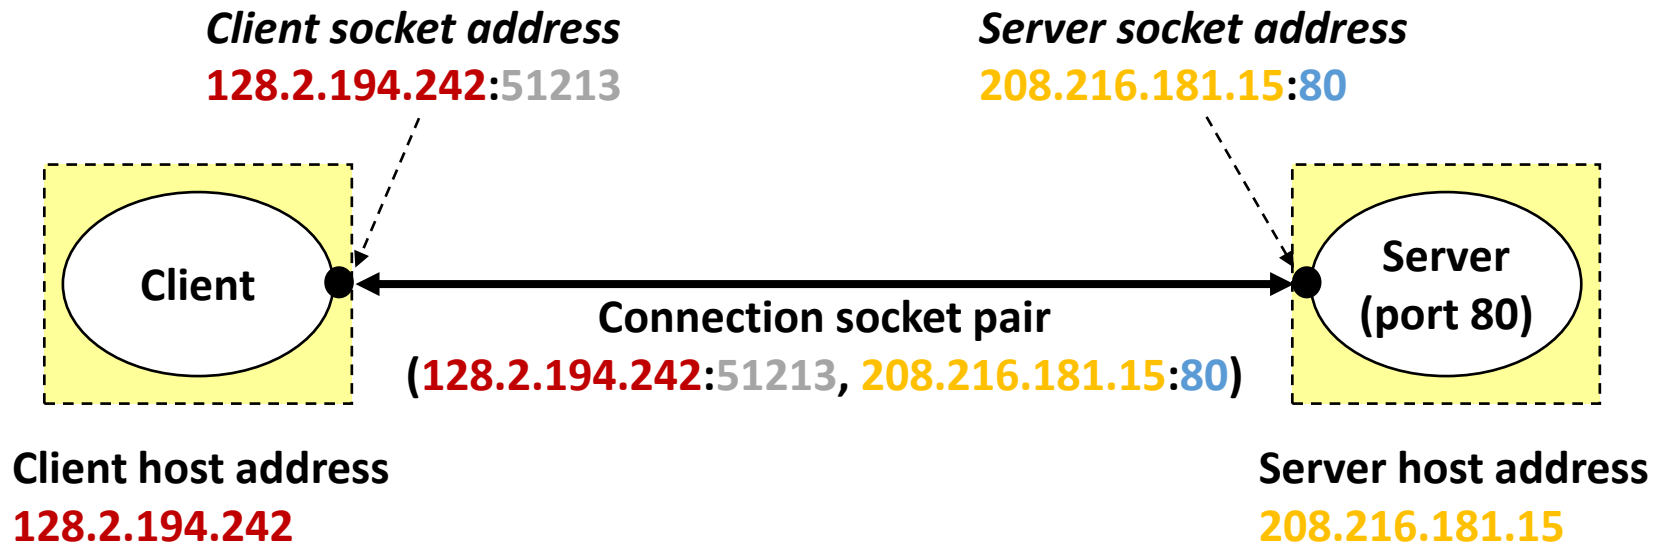
TA – Luke Albano ([lukealbano@arcs.skku.edu](mailto:lukealbano@arcs.skku.edu))

Sungkyunkwan University

# Internet Connection (1)

- Clients and servers communicate by sending streams of bytes over connections:
  - Point-to-point, full-duplex, and reliable
- A **socket** is an endpoint of a connection
  - Socket address is an <IP address:port number> pair
- A port is a 16-bit integer that identifies a process
  - **Ephemeral port**: assigned automatically on client when client makes a connection request
  - **Well-known port**: associated with some service provided by a server (e.g. port 80 is associated with web servers)
- A connection is uniquely identified by the socket addresses of its endpoints (**socket pair**)
  - <client IP:client port no., server IP:server port no.>

# Internet Connection (2)




*Note: 51213 is an ephemeral port allocated by the kernel*

*Note: 80 is a well-known port associated with Web servers*

# Client-Server Model

- Most network application is based on the client-server model
  - A **server** process and one or more **client** processes
    - Clients and servers are processes running on hosts (can be the same or different hosts)
  - Server manages some **resources**
  - Server provides **service** by manipulating resource for clients

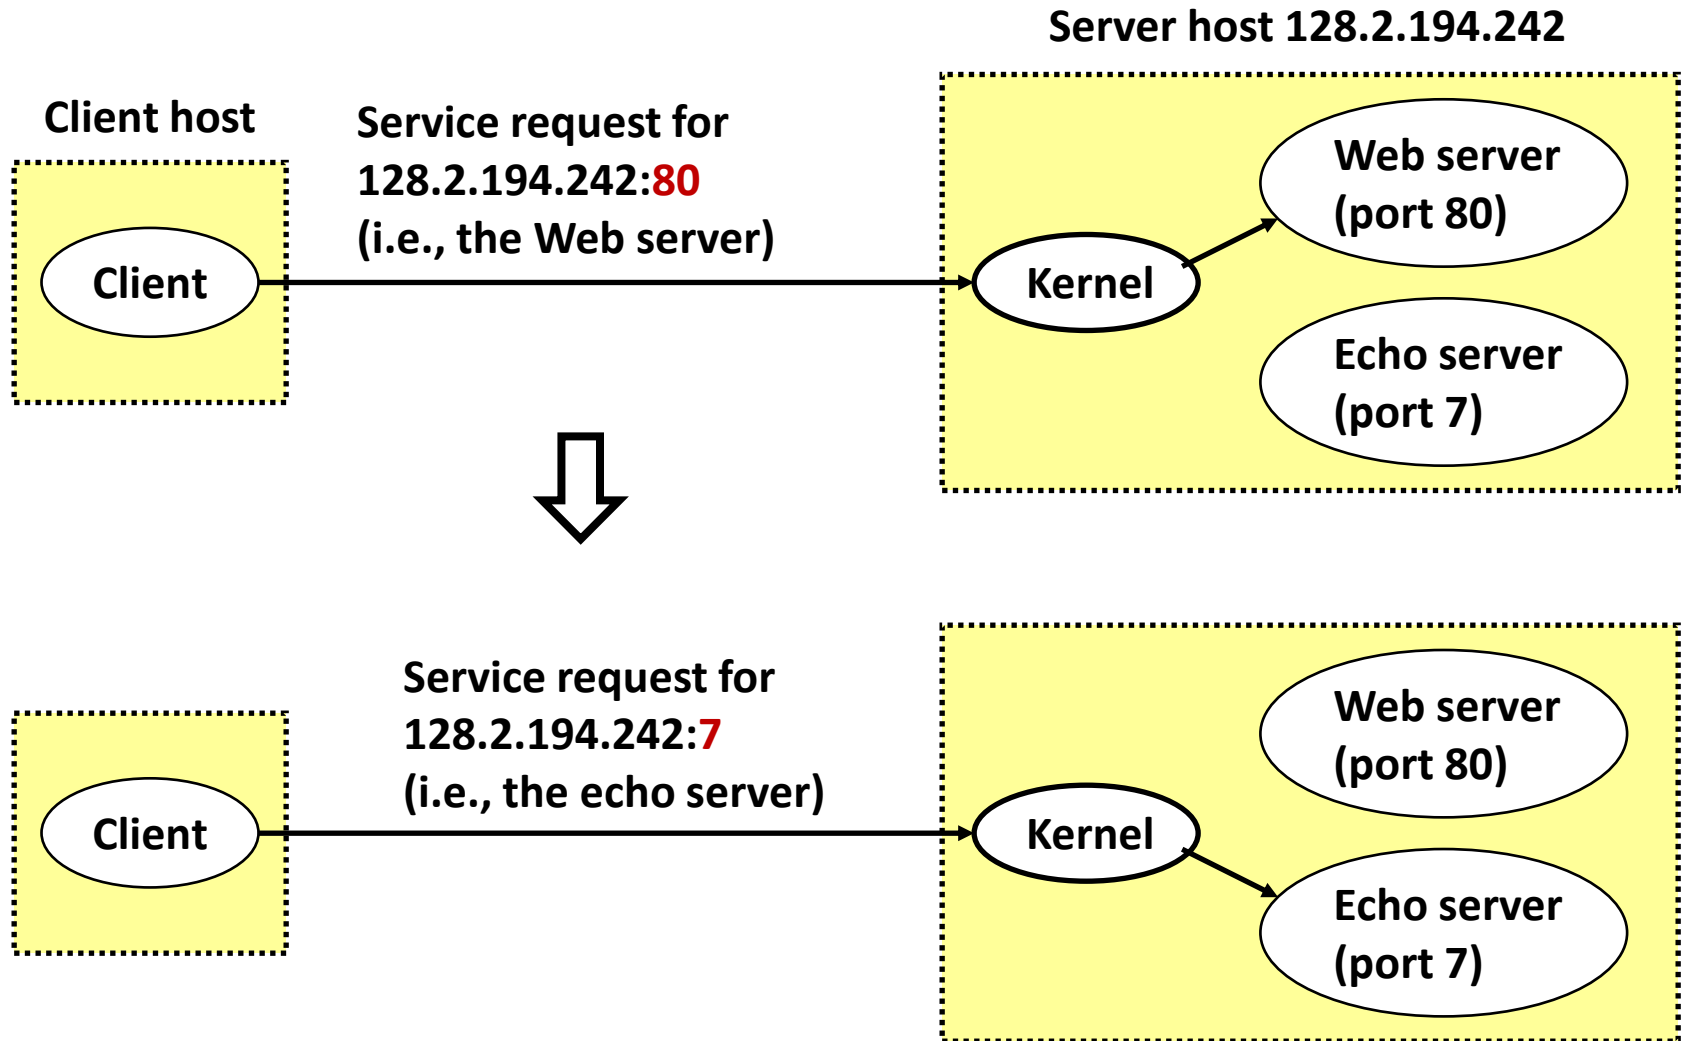


# Clients

- Examples of client programs
  - Web browsers, FTP, telnet, ssh, etc.
- How does a client find the server?
  - **IP address** in server socket address identifies the host
  - The **(well-known) port** in server socket address identifies the service; thus, port implicitly identifies the server process that performs the service
  - Examples of well-known ports (cf. **/etc/services**)
    - Port 21: FTP
    - Port 23: telnet
    - Port 25: mail
    - Port 80: web

```
ftp      21/tcp
fsp      21/udp      fspd
ssh      22/tcp          # SSH Remote Login Protocol
telnet   23/tcp
smtp     25/tcp      mail
time     37/tcp      timeserver
time     37/udp      timeserver
whois    43/tcp      nicname
finger   79/tcp
http     80/tcp      www      # WorldWideWeb HTTP
ntp      123/udp      # Network Time Protocol
```

# Using Ports



# Servers

- Servers are long-running processes (**daemons**)
  - Created at boot-time (typically) by the *init* process (process 1)
  - Run continuously until the machine is turned off
- Each server waits for requests to arrive on a well-known port associated with a particular service
  - e.g., port 21: FTP server, port 80: HTTP server
- A machine that runs a server process is often referred to as “server”

# Socket (1)

- Socket interface (circa 1981)
  - Was Introduced in BSD4.1 UNIX, 1981
  - Is based on **client-server paradigm**
  - Provides a user-level interface to the network
  - Can be explicitly created, used, released by applications
  - Consists of two types of transport service
    - Reliable, connection-oriented byte stream
    - Unreliable datagram
  - Is the underlying basis for all Internet applications



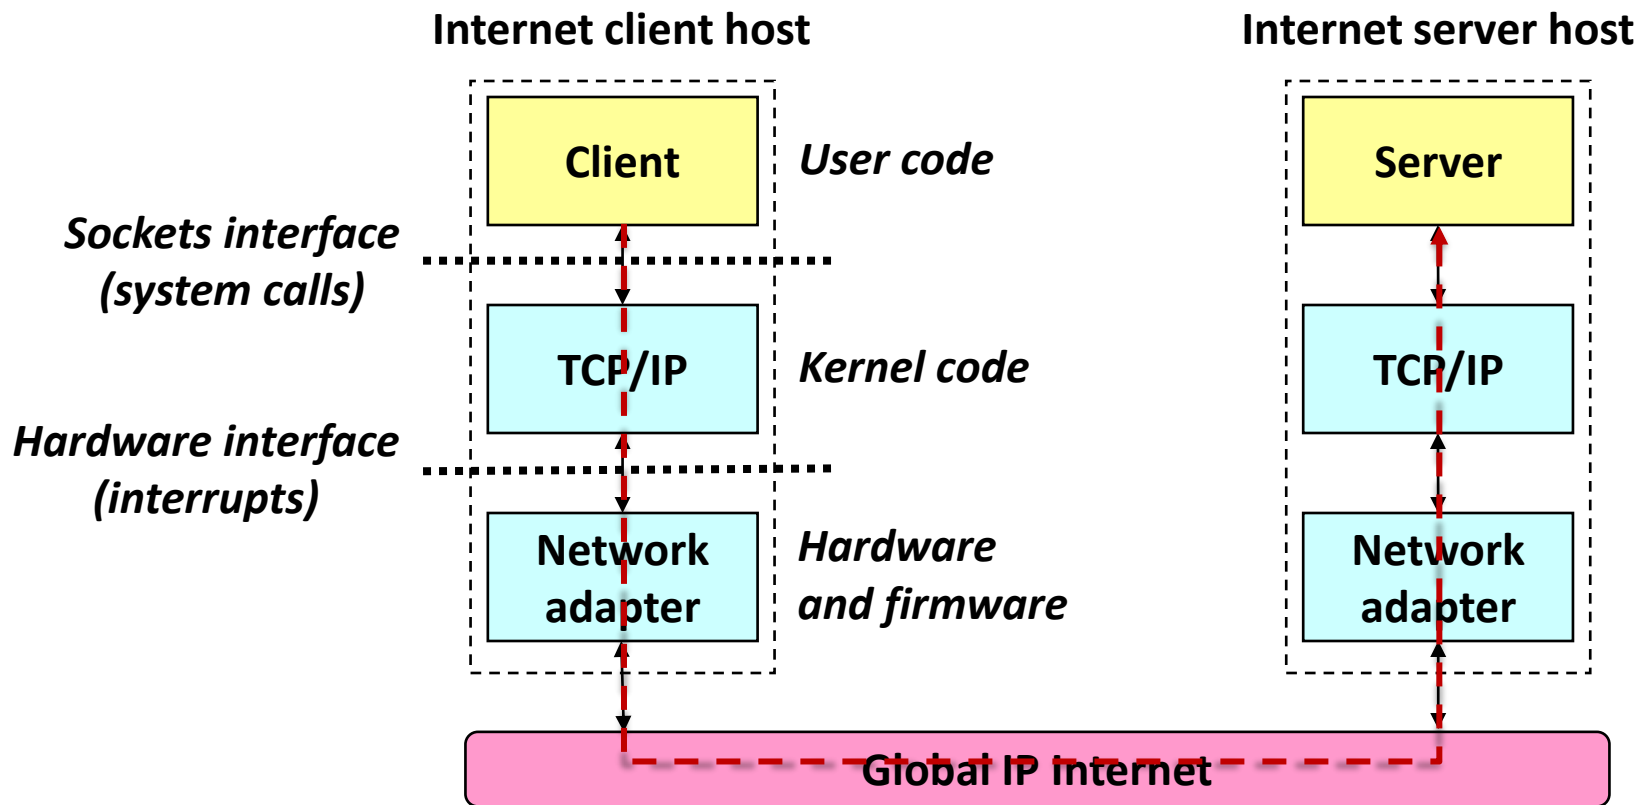
# Socket (2)

## ■ What is a **socket**?

- Remote-local, application-created/owned, OS-controlled interface to network (“door”)
  - To the kernel, socket is an endpoint of communication
  - To the application, socket is a file descriptor
    - Applications read/write from/to the network **using the file descriptor**
    - Note: All Unix I/O devices, including networks, are modeled **as files**
- Clients and servers communicate with each other by reading from and writing to socket descriptors
  - The main distinction between regular file I/O and socket I/O is how applications “open” socket descriptors

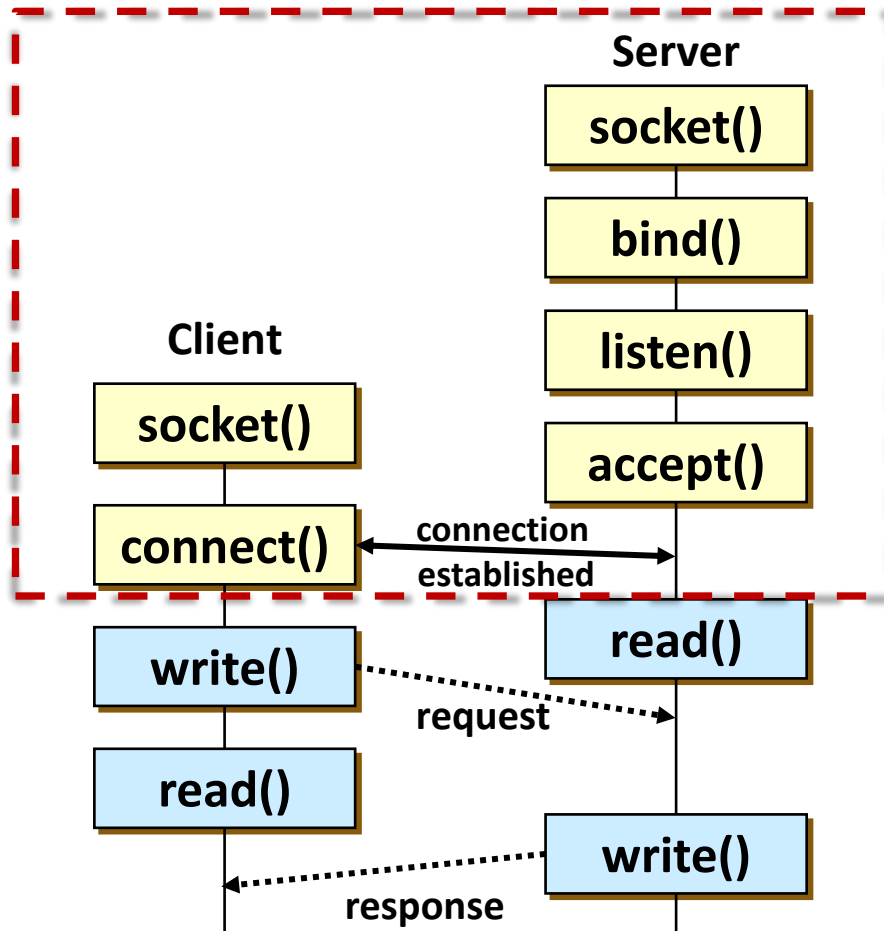
# Socket (3)

- Hardware/Software organization of an Internet application

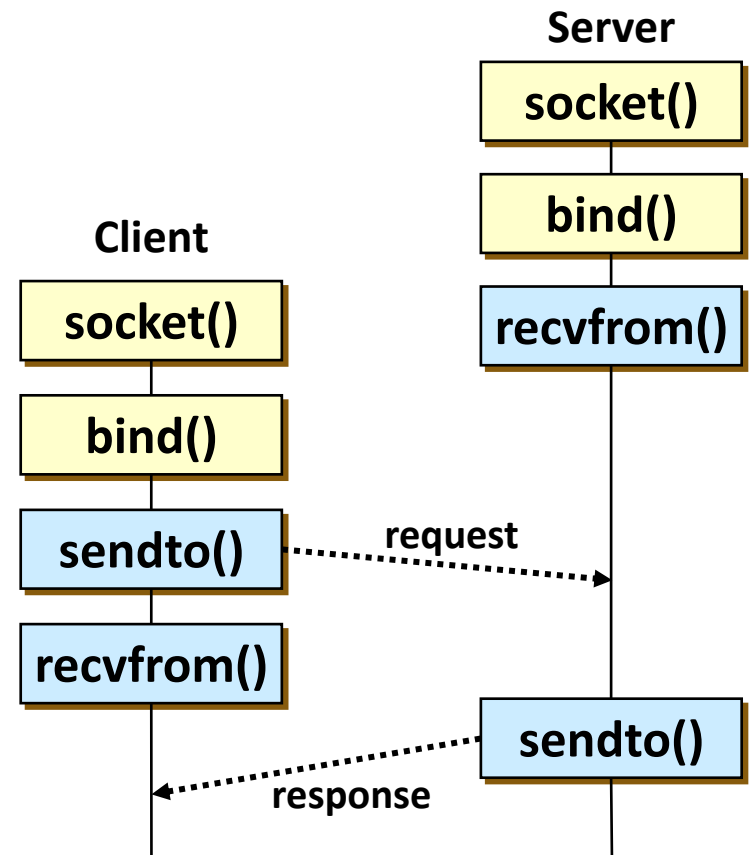


# TCP vs. UDP

## Connection-oriented service



## Connectionless service



# Socket Address Structure

- Generic socket address

- For address arguments to **connect()**, **bind()**, and **accept()**

```
struct sockaddr {  
    unsigned short  sa_family;    /* protocol family */  
    char            sa_data[14]; /* address data    */  
};
```

- Internet-specific socket address

- Must cast (**sockaddr\_in \***) to (**sockaddr \***)  
for **connect()**, **bind()**, and **accept()**

```
struct sockaddr_in {  
    unsigned short  sin_family; /* address family (always AF_INET) */  
    unsigned short  sin_port;   /* port num in network byte order */  
    struct in_addr   sin_addr;   /* IP addr in network byte order  */  
    unsigned char    sin_zero[8]; /* pad to sizeof(struct sockaddr) */  
};
```

# socket()

- int socket (int **domain**, int **type**, int **protocol**)
  - **socket()** creates a socket descriptor
  - **domain** specifies the communication domain
    - **AF\_UNIX**: Local Unix domain protocols
    - **AF\_INET**: IPv4 Internet protocols
  - **type** specifies the communication semantics
    - **SOCK\_STREAM**: provides *sequenced, reliable, two-way, connection-based* **byte streams**
    - **SOCK\_DGRAM**: supports **datagrams** (*connectionless, unreliable messages of a fixed maximum length*)
    - **SOCK\_RAW**: provides raw network protocol access
  - **protocol** specifies a particular protocol to be used with the socket (`IPPROTO_IP (0)`, `IPPROTO_TCP`, `IPPROTO_UDP`)

# connect()

- int connect (int **sockfd**,  
                  const struct sockaddr \***servaddr**,  
                  socklen\_t **addrlen**)
  - Used by a **TCP client** to establish a connection with a TCP server
  - **servaddr** contains **<IP address, port number>** of the server
  - Client does not have to call **bind()** before calling **connect()**
    - Kernel will choose both ephemeral port and source IP address if necessary
  - Client process **suspends (blocks)** until the connection is created

# Functions for Echo Client

MEMSET(3)

Linux Programmer's Manual

MEMSET(3)

## NAME

memset - fill memory with a constant byte

## SYNOPSIS

```
#include <string.h>
```

```
void *memset(void *s, int c, size_t n);
```

## DESCRIPTION

The `memset()` function fills the first `n` bytes of the memory area pointed to by `s` with the constant byte `c`.

## RETURN VALUE

The `memset()` function returns a pointer to the memory area `s`.

MEMCPY(3)

Linux Programmer's Manual

MEMCPY(3)

## NAME

memcpy - copy memory area

## SYNOPSIS

```
#include <string.h>
```

```
void *memcpy(void *dest, const void *src, size_t n);
```

## DESCRIPTION

The `memcpy()` function copies `n` bytes from memory area `src` to memory area `dest`. The memory areas must not overlap. Use `memmove(3)` if the memory areas do overlap.

## RETURN VALUE

The `memcpy()` function returns a pointer to `dest`.

# Echo Client (1)

```
// swe2024-41_23s/2023s/w10/echo_client.c
```

```
#define MAXLINE 80
```

```
int main (int argc, char *argv[]) {
    ssize_t num_bytes; // number of bytes read
    int num_bytes, socket_fd; // fd that handles socket
    struct hostent *host_entry; // contains info related to host
    struct sockaddr_in socket_address;
    char buffer[MAXLINE];

    if (argc < 3){
        printf("Received %d arguments. Please enter hostname and port
number!\n", argc - 1);
        exit(1);
    }

    char *host = argv[1]; // contains hostname
    uint32_t port = (uint32_t)strtol(argv[2], NULL, 10); // contains
port number
```

```
#include <arpa/inet.h> // htons, tcp, ipv4
#include <sys/types.h> // ssize_t
#include <sys/socket.h> // socket
#include <netdb.h> // hostent
#include <stdio.h> // printf
#include <unistd.h> // write, read
#include <stdlib.h> // atoi, exit
#include <string.h> // memset, memcpy
```



# Echo Client (2)

```
// creates a socket for TCP connections
// AF_INET: IPv4, SOCK_STREAM: stream socket, IPPROTO_TCP: TCP
if ((socket_fd = socket(AF_INET, SOCK_STREAM, IPPROTO_TCP)) < 0) {
    printf("socket() failed.\n");
    exit(2);
}

// Gets the host entry given the hostname
if ((host_entry = gethostbyname(host)) == NULL) {
    printf("invalid hostname %s\n", host);
    exit(3);
}

memset((char *)&socket_address, 0, sizeof(socket_address)); //
initializes socket_address fields to zero
socket_address.sin_family = AF_INET; // sets address family to IPv4
memcpy((char *)&socket_address.sin_addr.s_addr, (char *)host_entry-
>h_addr, host_entry->h_length); // copies ip address
socket_address.sin_port = htons(port); // sets port after converting
it to "network byte order"
```

# Echo Client (3)

```
// Connects the socket to the given IP address
if (connect(socket_fd, (struct sockaddr
*)&socket_address, sizeof(socket_address)) < 0) {
    printf("connect() failed.\n");
    exit(4);
}

// Reads user input
while ((num_bytes = read(STDIN_FILENO, buffer, MAXLINE)) >
0) {
    write(socket_fd, buffer, num_bytes); // Send input to
server
    num_bytes = read(socket_fd, buffer, MAXLINE); // Get
response from server
    write(STDOUT_FILENO, buffer, num_bytes); // Prints
response to server
}
close(socket_fd); // Destroys socket
}
```

# Example Output

```
nvim echo_client.c && gcc echo_client.c -Wall -Wextra -o echo_client && ./echo_client tcpbin.com 4242
```

Hi!

Hi!



안녕하세요~

안녕하세요~

# bind()

- int bind (int **sockfd**,  
const struct sockaddr \***addr**,  
socklen\_t **addrlen**)
  - **bind()** gives socket **sockfd** the local address **addr**
  - **addr** is **addrlen** bytes long
  - Servers bind their well-known port when they start
  - If a TCP server binds a specific IP address to its socket, this **restricts** the socket to receive incoming client connections **destined only to that IP address**
  - Normally, a TCP client lets the kernel choose the ephemeral port and client IP address

# listen()

- int listen (int **sockfd**, int **backlog**)
  - **listen()** converts an unconnected socket into a **passive socket**, indicating that the **kernel should accept incoming connection requests**
    - When socket is created, it is assumed to be an active socket, that is, a client socket that will issue **connect()**
  - **backlog** specifies the **maximum number of connections** that kernel should queue for this socket
  - Historically, a backlog of 5 was used, as that was maximum value supported by 4.2BSD
    - Busy HTTP servers must specify much larger backlog, and newer kernels must support larger values


# accept() (1)

- int accept (int **sockfd**,  
              struct sockaddr \***cliaddr**,  
              socklen\_t \***addrlen**)
  - **accept()** **blocks waiting** for a connection request
  - **accept()** returns a **connected descriptor** with the same properties as the **listening descriptor**
    - Kernel creates one connected socket for each client connection that is accepted
    - **accept()** returns when **the connection** between client and server is **created and ready for I/O transfers**
    - All I/O with the client will be done via the connected socket
  - The **cliaddr** and **addrlen** arguments are used to return the address of the **connected peer process** (the client)


# accept() (2)



**1. Server blocks in `accept`, waiting for connection request on listening descriptor `listenfd`**



**2. Client makes connection request by calling and blocking in `connect`**



**3. Server returns `connfd` from `accept`. Client returns from `connect`. Connection is now established between `clientfd` and `connfd`.**

# accept() (3)

- Listening descriptor
  - End point for client connection requests
  - Created once and **exists for lifetime of the server**
- Connected descriptor
  - End point of the **connection between client and server**
  - **A new descriptor** is created **each** time the server accepts a connection request from a **client**
  - It exist only as long as it takes to service client
- Why the distinction?
  - Allow for **concurrent servers** that can communicate over many client connections simultaneously



# Echo Server (1)

```
#include <arpa/inet.h> // htons/l, tcp, ipv4
#include <sys/types.h> // ssize_t
#include <sys/socket.h> // socket
#include <stdio.h> // printf
#include <unistd.h> // write, read
#include <stdlib.h> // atoi, exit
#include <string.h> // memset, memcpy
```

```
// swe2024-41_23s/2023s/w10/echo_server.c
```

```
#define MAXLINE 80
```

```
int main (int argc, char *argv[]) {
    ssize_t num_bytes; // number of bytes read
    size_t connected_address_length;
    int listening_fd, connected_fd;
    struct sockaddr_in socket_address, connected_address;
    char buffer[MAXLINE];

    if (argc < 2){
        printf("Received %d arguments. Please enter port number!\n", argc - 1);
        exit(1);
    }

    uint32_t port = (uint32_t)strtol(argv[1], NULL, 10); // contains port
    number

    // creates a socket for TCP connections
    // AF_INET: IPv4, SOCK_STREAM: stream socket, IPPROTO_TCP: TCP
    if ((listening_fd = socket(AF_INET, SOCK_STREAM, IPPROTO_TCP)) < 0) {
        printf("socket() failed.\n");
        exit(2);
    }
}
```

# Echo Server (2)

```
memset((char *)&socket_address, 0, sizeof(socket_address)); //
initializes socket_address fields to zero
socket_address.sin_family = AF_INET; // sets address family to
IPv4
socket_address.sin_addr.s_addr = htonl(INADDR_ANY); // copies
ip address
socket_address.sin_port = htons(port); // sets port after
converting it to "network byte order"

// Assigns socket address to the listening fd
if (bind(listening_fd, (struct sockaddr *)&socket_address,
sizeof(socket_address)) < 0) {
    printf("bind() failed.\n");
    exit(3);
}

// Listens to 5 connections at most
if (listen(listening_fd, 5) < 0) {
    printf("listen() failed.\n");
    exit(4);
}
```

# Echo Server (3)

```
while (1) {
    connected_address_length = sizeof(connected_address);


    // Accepts connection
    if ((connected_fd = accept(
        listening_fd,
        (struct sockaddr *)&connected_address,
        (socklen_t *)&connected_address_length
    )) < 0) {
        printf("accept() failed.\n");
        continue;
    }


    // Receives data from client
    while ((num_bytes = read(connected_fd, buffer, MAXLINE)) > 0) {
        printf("got %ld bytes from client.\n", num_bytes);
        write(connected_fd, buffer, num_bytes);
    }

    // Terminates connection
    printf("connection terminated.\n");
    close(connected_fd);
}
```

# Example Output

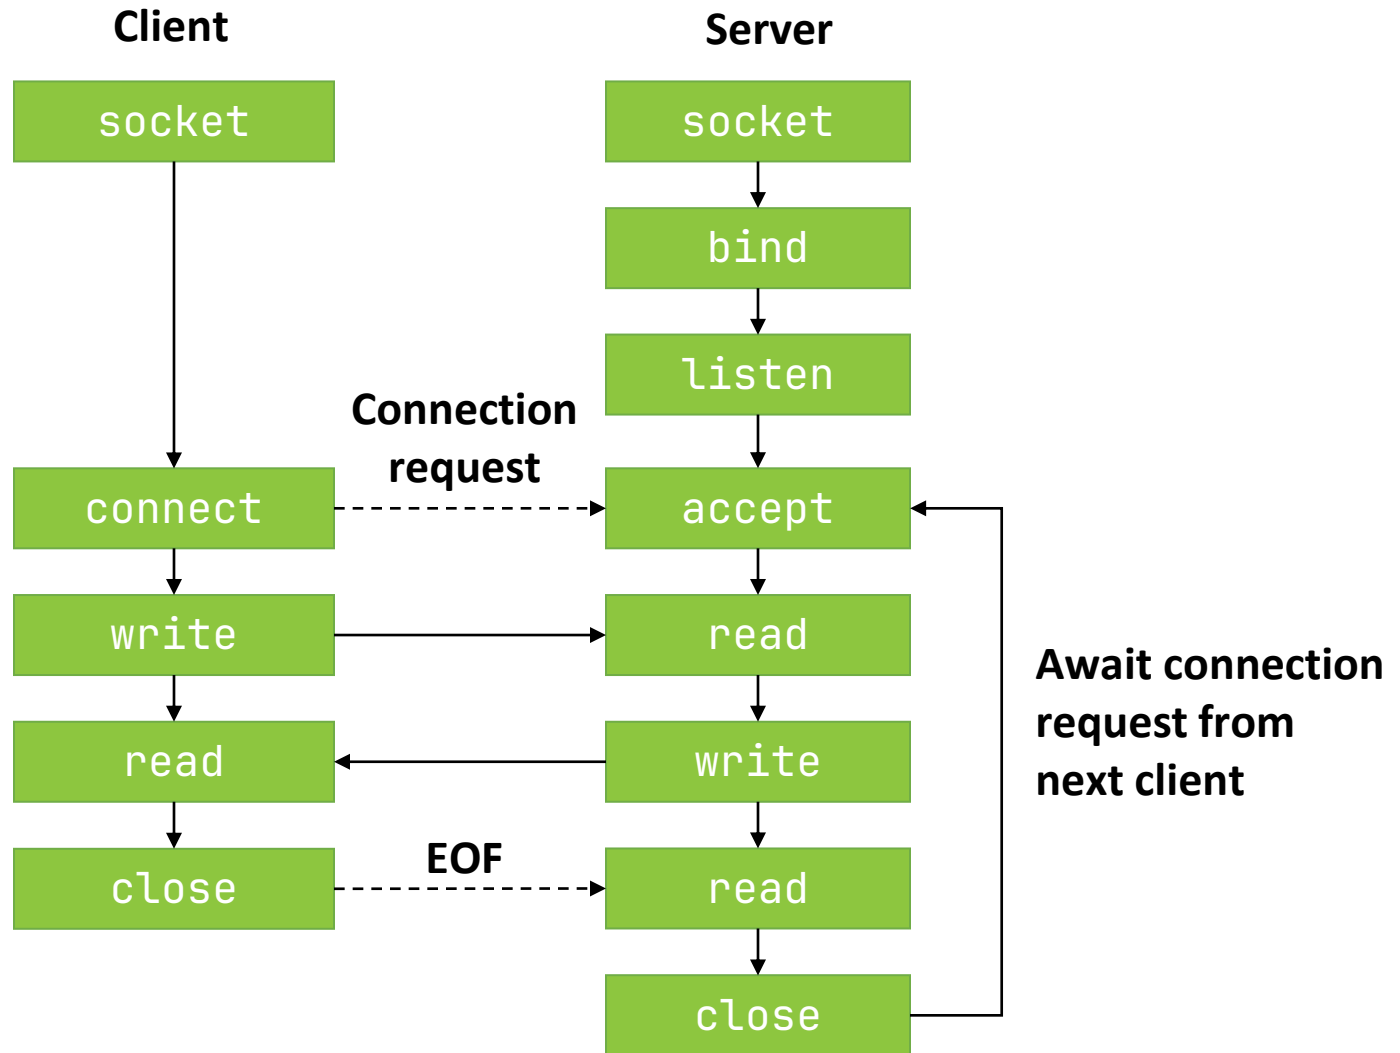
```

 ./echo_server 202441 (base)
got 4 bytes from client.
got 17 bytes from client.
got 5 bytes from client.

 ./echo_client localhost 202441 (base)
Hi!
Hi!
안녕하세요~
안녕하세요~
😊
😊

```

# Echo Server (4)



# Lab Exercise

## ■ Make *file transfer*

- Server runs in background by attaching & when running
- Client runs in foreground and is shown on the screen
- Client should be under an infinite loop with **conditional exit** (“quit”)
- If server receives a ‘quit’ string, the server shuts down
- When name of the file is entered, client reads file and sends it to server
- Server saves file name received from client by appending *\_copy*
- File is text file and does not have extension
  - e.g., a.txt(X), b.exe(X), c.c(X), Aladdin(O), a(O), b(O)
- Size of file is smaller than 200,000 bytes
- Length of file name is shorter than 50 bytes
- When testing, use last 5 digits of your student ID for the port

# Lab Exercise

- Skeleton code

```
cp ~swe2024-41_23s/2023s/p10 ./ -r
```

- When testing on the cluster, use last 5 digits of your student ID for the port (i.e., do not use 10080!)

- Example)

```
root@ubuntu:/home/seungwoo/swe2033/ex10# ls
a  b  client  server
root@ubuntu:/home/seungwoo/swe2033/ex10# ./server 10080 &
[1] 54822
root@ubuntu:/home/seungwoo/swe2033/ex10# ./client localhost 10080
a
got 31 bytes from client.
Send 31 bytes to server.
b
got 17 bytes from client.
Send 17 bytes to server.
quit
[1]+  완료                  ./server 10080
root@ubuntu:/home/seungwoo/swe2033/ex10# ls
a  a_copy  b  b_copy  client  server
```

# Exercise Submission

- Submit your exercise code and Makefile
  - InUiYeJi cluster
  - Submit the folder into p10
  - ~swe2024-41\_23s/bin/submit p10 p10
  - Due date: Friday, 05 May 2023, 23:59
  - We will compile by using command *make*
    - If compilation fails, your points for this exercise will be **zero**

```
./p10
```

```
server.c  
client.c  
Makefile
```



# Summary Report

- Summary report about man command result of
  - socket()
  - bind()
  - listen()
  - accept()
  - connect()
- Submission form
  - A4 size PDF format (No page limitation)
  - [SWE2024 Report-10] studentID\_name
  - Ex) [SWE2024 Report-10] 2022XXXXXX\_홍길동
  - Submit to iCampus
  - Due by Friday, 5 May 2023, 23:59