

Pthreads

Prof. Joonwon Lee (joonwon@skku.edu)

TA – Jaehyung Park (jaeseanpark@gmail.com)


TA – Luke Albano (lukealbano@arcs.skku.edu)

Sungkyunkwan University

Thread vs. Process


■ Process

- Process has its own address space
- Each process has its own data
 - e.g., global variables, stack, heap



■ Thread

- Multiple thread share on address space
 - But its own **stack** and **register** context
- Threads within same address space share data
 - e.g., global variables, heap



POSIX Threads

- Pthreads = POSIX thread
- Standard thread APIs for UNIX
 - IEEE standard 1003.1c-1995
 - Allow program to control multiple threads that overlap in time

Pthreads APIs

- Thread management APIs
 - Create, terminate, and join threads
 - Set and query thread attributes
- Mutex APIs
 - Create, destroy, lock, and unlock mutexes
- Conditional variable APIs
 - Create and destroy condition variables
 - Wait and signal based upon specified variable values

Pthreads Usage

- Pthreads header

- Header file (*pthread.h*) is required to use Pthreads APIs
include <pthread.h>

- Compilation

- Almost all compilers support compilation with Pthreads library
- In GNU compiler, add *-lpthread* to compile Pthreads programs
\$ gcc -Wall -o hello hello.c -lpthread

Data Types of Pthreads

- **pthread[_object]_t**
 - **pthread_t**
 - Handle of a thread
 - Contain thread ID after creating thread
 - **pthread_attr_t**
 - Handle thread attribute type
 - Set attribute of newly created thread
 - **pthread_mutex_t, pthread_mutexattr_t**
 - **pthread_cond_t, pthread_condattr_t**
 - etc.

Thread Creation

- **int pthread_create (pthread_t *thread,
pthread_attr_t *attr,
void *(*start_routine)(void*),
void *arg)**
 - Return new **thread ID** via *thread* parameter
 - Thread ID is useful for various operations on thread
 - *attr* parameter is used to set thread attributes
 - NULL for default values
 - *start_routine* parameter denotes the **C routine**
 - Thread will execute once it is created
 - A single argument may be passed to *start_routine* via *arg*
 - For multiple *args*, store them in a structure and pass its pointer
 - Return 0, on successful thread creation
 - Non-zero if unsuccessful

Thread Termination

- **void pthread_exit (void *retval)**
 - Terminate execution of calling thread
 - After a thread has completed its work typically
 - A thread is no longer required to exist
 - *retval* parameter is return value of thread
 - This value can be consulted from other threads with **pthread_join()**
 - It does not close files
 - Any files opened inside the thread will remain open after thread is terminated

Thread Creation & Termination Example

```
#define NUM_THREADS 4

void *thread(void *arg) {
    long id = (long)arg;
    printf("thread#%ld: Hello Thread!\n", id);
    pthread_exit(NULL);
}

int main() {
    pthread_t tid[NUM_THREADS];
    long t;

    for (t = 0; t < NUM_THREADS; t++) {
        printf("main: creating thread#%ld\n", t);
        if (pthread_create(&tid[t], NULL, thread, (void*)t)) {
            printf("ERROR: pthread creation failed.\n");
            exit(1);
        }
    }
    printf("main: bye bye!\n");
    pthread_exit(NULL);
}
```

Thread Joining

- **int pthread_join (pthread_t thread, void **retval)**
 - Suspend execution of caller thread
 - Until the thread identified by *thread* parameter terminates
 - Thread terminates by
 - ✓ Calling **pthread_exit()**
 - ✓ Cancelled from other thread (such as **pthread_cancel()**)
 - Return value is stored in the location pointed by *retval* parameter
 - **PTHREAD_CANCELLED** is stored if thread was cancelled
 - Thread must be joinable
 - It is impossible to join a detached thread
 - Returns 0 on success
 - Non-zero if unsuccessful

Thread Join Example

```
#define NUM_THREADS 4

void *thread(void *arg) {
    long id = (long)arg;
    printf("thread#%ld: Hello Thread!\n", id);
    pthread_exit(NULL);
}

int main() {
    pthread_t tid[NUM_THREADS];
    long t;

    for (t = 0; t < NUM_THREADS; t++) {
        printf("main: creating thread#%ld\n", t);
        if (pthread_create(&tid[t], NULL, thread, (void*)t)) {
            printf("ERROR: pthread creation failed.\n");
            exit(1);
        }
    }
    for (t = 0; t < NUM_THREADS; t++) {
        pthread_join(tid[t], NULL);
    }
    printf("main: bye bye!\n");
    return 0;
}
```

Thread Detachment

- **int pthread_detach (pthread_t thread)**
 - Put thread in detached state
 - This guarantees that memory resources consumed by *thread* parameter will be freed immediately when thread terminates
 - However, this prevents other threads from synchronizing on termination of thread by calling **pthread_join()**
 - Return 0 on success
 - Thread can be detached when it is created:

```
pthread_t tid;
pthread_attr_t attr;

pthread_attr_init (&attr);
pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_DETACHED);
pthread_create(&tid, &attr, start_routine, NULL);
pthread_attr_destroy (&attr);
```

Thread Cancellation

- **int pthread_cancel (pthread_t thread)**
 - Send a cancellation request to a thread
 - This thread has same ID to *thread* parameter
 - Status of target thread determines the cancellation
 - **PTHREAD_CANCEL_ENABLE**: allow the cancellation
 - **PTHREAD_CANCEL_DISABLE**: disallow the cancellation; cancellation request is blocked until the thread can be cancelled
 - Type of target thread determines the reaction
 - **PTHREAD_CANCEL_ASYNCHRONOUS**: thread can be cancelled at anytime
 - **PTHREAD_CANCEL_DEFERRED**: cancellation request is deferred until the cancellation point
 - Return 0 if the cancellation successful
 - Non-zero if unsuccessful

Thread Identifiers

- **pthread_t pthread_self (void)**
 - Return unique, system assigned thread ID of calling thread
- **int pthread_equal (pthread_t t1, pthread_t t2)**
 - Return non-zero value if parameter t1 & t2 refer to same thread
 - C language equivalence operator (==) should not be used to compare two thread IDs against each other
 - This is because thread IDs are opaque objects
 - ✓ Pthreads library identifies a unique thread using this object

Thread Attributes

- **pthread_attr_t**
- **int pthread_attr_init (pthread_attr_t *attr)**
 - Initialize attribute objects to their default value
- **int pthread_attr_destroy (pthread_attr_t *attr)**
 - Delete attribute objects initialized by **pthread_attr_init()**
 - Must be re-initialized to use destroyed object again

Thread Attributes (Cont.)

- **int pthread_attr_setdetachstate (pthread_attr_t *attr, int detachstate)**
 - Set attribute of thread as either detached or joinable
 - Value in *detachstate* parameter is applied to detach state attribute
 - e.g., PTHREAD_CREATE_DETACHED, PTHREAD_CREATE_JOINABLE
- **int pthread_attr_getdetachstate (pthread_attr_t *attr, int *detachstate)**
 - Retrieve thread attribute to find it is detached or joinable
 - Thread attribute state is returned in *detachstate* parameter

Thread-Local Storage (TLS)

- TLS is a computer programming method
 - It uses static or global memory local to a thread
- Efficient and flexible usage of thread-local data
- TLS in C language
 - Cooperation between compiler/linker and runtime system
 - In GNU C, use the new keyword `__thread`
 - Global variable – `__thread int number;`
 - Static variable – `static __thread int number;`

Thread-Local Storage Example

```
#define NUM_THREADS 16

int value;

void *thread(void *arg) {
    long id = (long)arg;
    value++;
    printf("thread#%ld - value: %d\n", id, value);
    pthread_exit(NULL);
}

int main() {
    pthread_t tid[NUM_THREADS];
    long t;
    for (t = 0; t < NUM_THREADS; t++) {
        if (pthread_create(&tid[t], NULL, thread, (void*)t)) {
            printf("ERROR: pthread creation failed.\n");
            exit(1);
        }
    }
    for (t = 0; t < NUM_THREADS; t++) {
        pthread_join(tid[t], NULL);
    }
    return 0;
}
```

```
jaehyun@csl:~/SSE/week11$ ./thread_local_storage
thread#0 - value: 1
thread#1 - value: 2
thread#2 - value: 3
thread#3 - value: 4
jaehyun@csl:~/SSE/week11$
```

Thread-Local Storage Example

```
#define NUM_THREADS 16
```

```
__thread int value;
```

```
void *thread(void *arg) {  
    long id = (long)arg;  
    value++;  
    printf("thread#%ld - value: %d\n", id, value);  
    pthread_exit(NULL);  
}
```

```
int main() {  
    pthread_t tid[NUM_THREADS];  
    long t;  
    for (t = 0; t < NUM_THREADS; t++) {  
        if (pthread_create(&tid[t], NULL, thread, (void*)t)) {  
            printf("ERROR: pthread creation failed.\n");  
            exit(1);  
        }  
    }  
    for (t = 0; t < NUM_THREADS; t++) {  
        pthread_join(tid[t], NULL);  
    }  
    return 0;  
}
```

```
jaehyun@cs1:~/SSE/week11$ ./thread_local_storage  
thread#0 - value: 1  
thread#3 - value: 1  
thread#2 - value: 1  
thread#1 - value: 1  
jaehyun@cs1:~/SSE/week11$
```

Exercise

- Matrix-vector multiplication using **multi-threading**
 - Multiplication result of **$M \times N$ matrix** and **$N \times 1$ vector** is **$M \times 1$ vector**
 - Row size (M), column size (N) of matrix are received by arguments
 - Elements of matrix and vector are **randomly** assigned between **0~9** (data type is **int**)
 - Create threads as much as the row size (M)
 - Each thread performs a calculation on one row of the matrix

Exercise

- Matrix-vector multiplication

$$\begin{array}{lcl} \text{Thread \#0} & \rightarrow & \\ \text{Thread \#1} & \rightarrow & \\ \text{Thread \#2} & \rightarrow & \end{array} \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} ax + by + cz \\ dx + ey + fz \\ gx + hy + iz \end{bmatrix}$$

- Restrictions

- Thread ID and calculation results are delivered through the *thread_data* struct
 - Do not use global variable for thread ID and calculation result
- Main function must wait other threads to terminate
 - Use `pthread_join()` system call

Exercise

- Random variable
 - Should include *time.h* header
 - `srand(time(NULL));`
 - `int a = rand() % 10;`

- Running example



```
jaehyun@cs1:~/SSE/week11/exercise$ gcc -o w11 exercise.c -lpthread
jaehyun@cs1:~/SSE/week11/exercise$ ls
exercise.c  skeleton.c  w11
jaehyun@cs1:~/SSE/week11/exercise$ ./w11 4 5
*** Matrix ***
[ 8 ] [ 0 ] [ 7 ] [ 6 ] [ 4 ]
[ 1 ] [ 3 ] [ 0 ] [ 4 ] [ 0 ]
[ 5 ] [ 7 ] [ 5 ] [ 1 ] [ 3 ]
[ 9 ] [ 7 ] [ 2 ] [ 9 ] [ 3 ]
*** Vector ***
[ 4 ]
[ 9 ]
[ 8 ]
[ 6 ]
[ 0 ]

*** Result ***
[ 124 ]
[ 55 ]
[ 129 ]
[ 169 ]
jaehyun@cs1:~/SSE/week11/exercise$
```

Exercise

- Submit your exercise source code
 - To InUiYeJi Cluster
 - Put your Makefile and *.c files in p11 folder
 - Submit using

```
$ ~swe2024-41_23s/bin/submit p11 p11
```
 - We will compile by using command *make*
 - When compilation fails, you get zero points
 - Compiled binary name should be “*p11*”
- Due 2023/5/12 23:59