

Concurrent Programming

Prof. Joonwon Lee (joonwon@skku.edu)

TA – Jaehyung Park (jaeseanpark@gmail.com)

TA – Luke Albano (lukealbano@arcs.skku.edu)

Sungkyunkwan University

Iterative Server

- Iterative echo server from previous lecture

```
int main (int argc, char *argv[]) {
    . . .

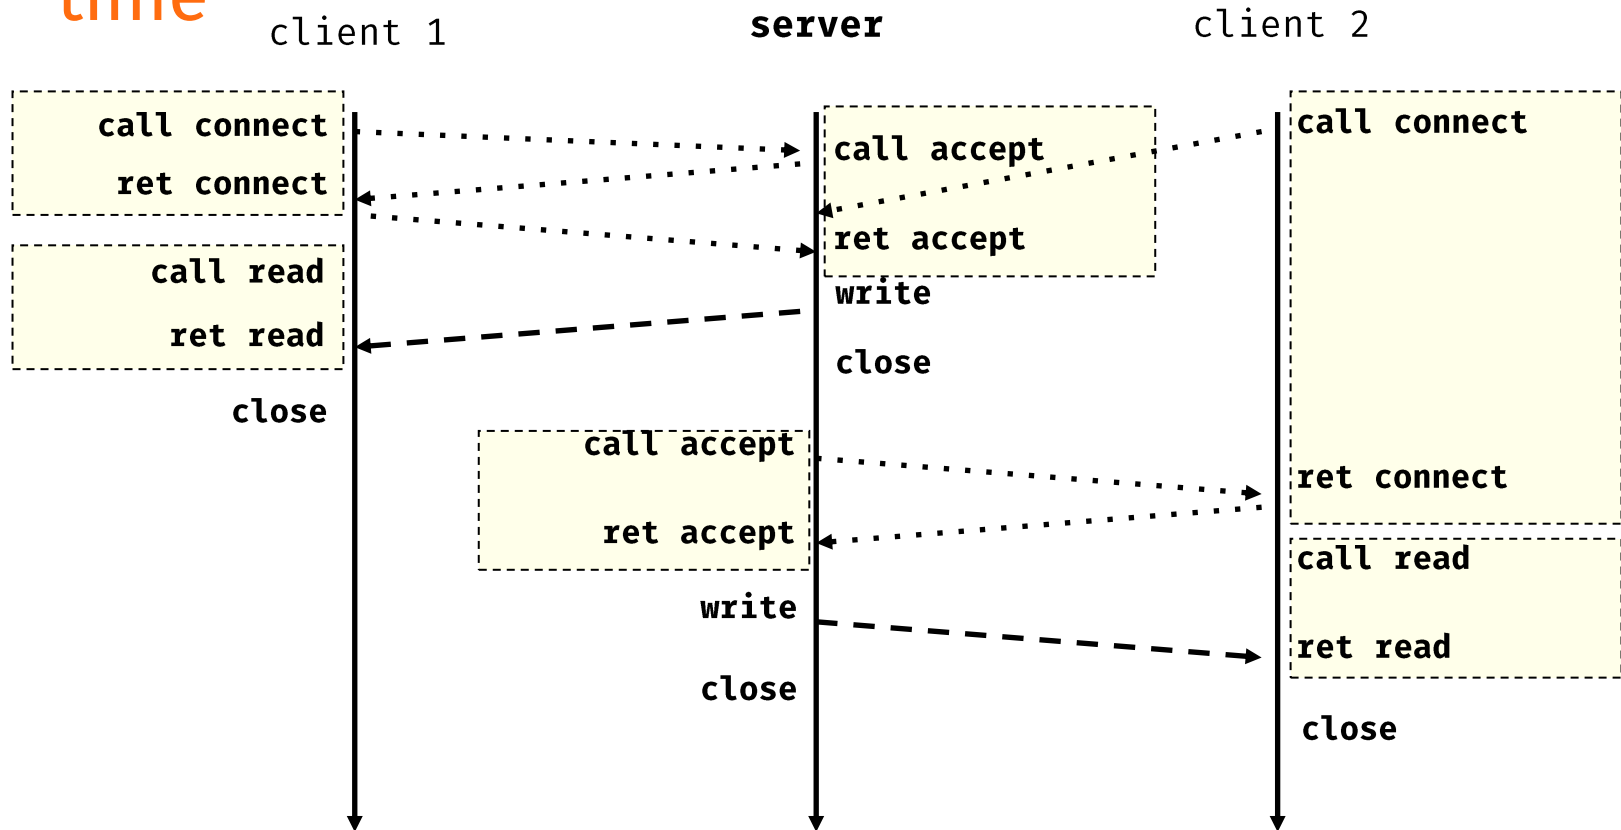
    while (1) {
        if ((connfd = accept(listenfd, (struct sockaddr *)&caddr,
                             (socklen_t *)&caddrlen)) < 0) {
            printf("accept() failed\n");
            continue;
        }

        while ((n = read(connfd, buf, MAXLINE)) > 0) {
            printf("got %d bytes from client.\n", n);
            write(connfd, buf, n);
        }


        close(connfd);
    }
}
```

Iterative Server

- Iterative server processes **one request at a time**



Fundamental Flaw of Iterative Server



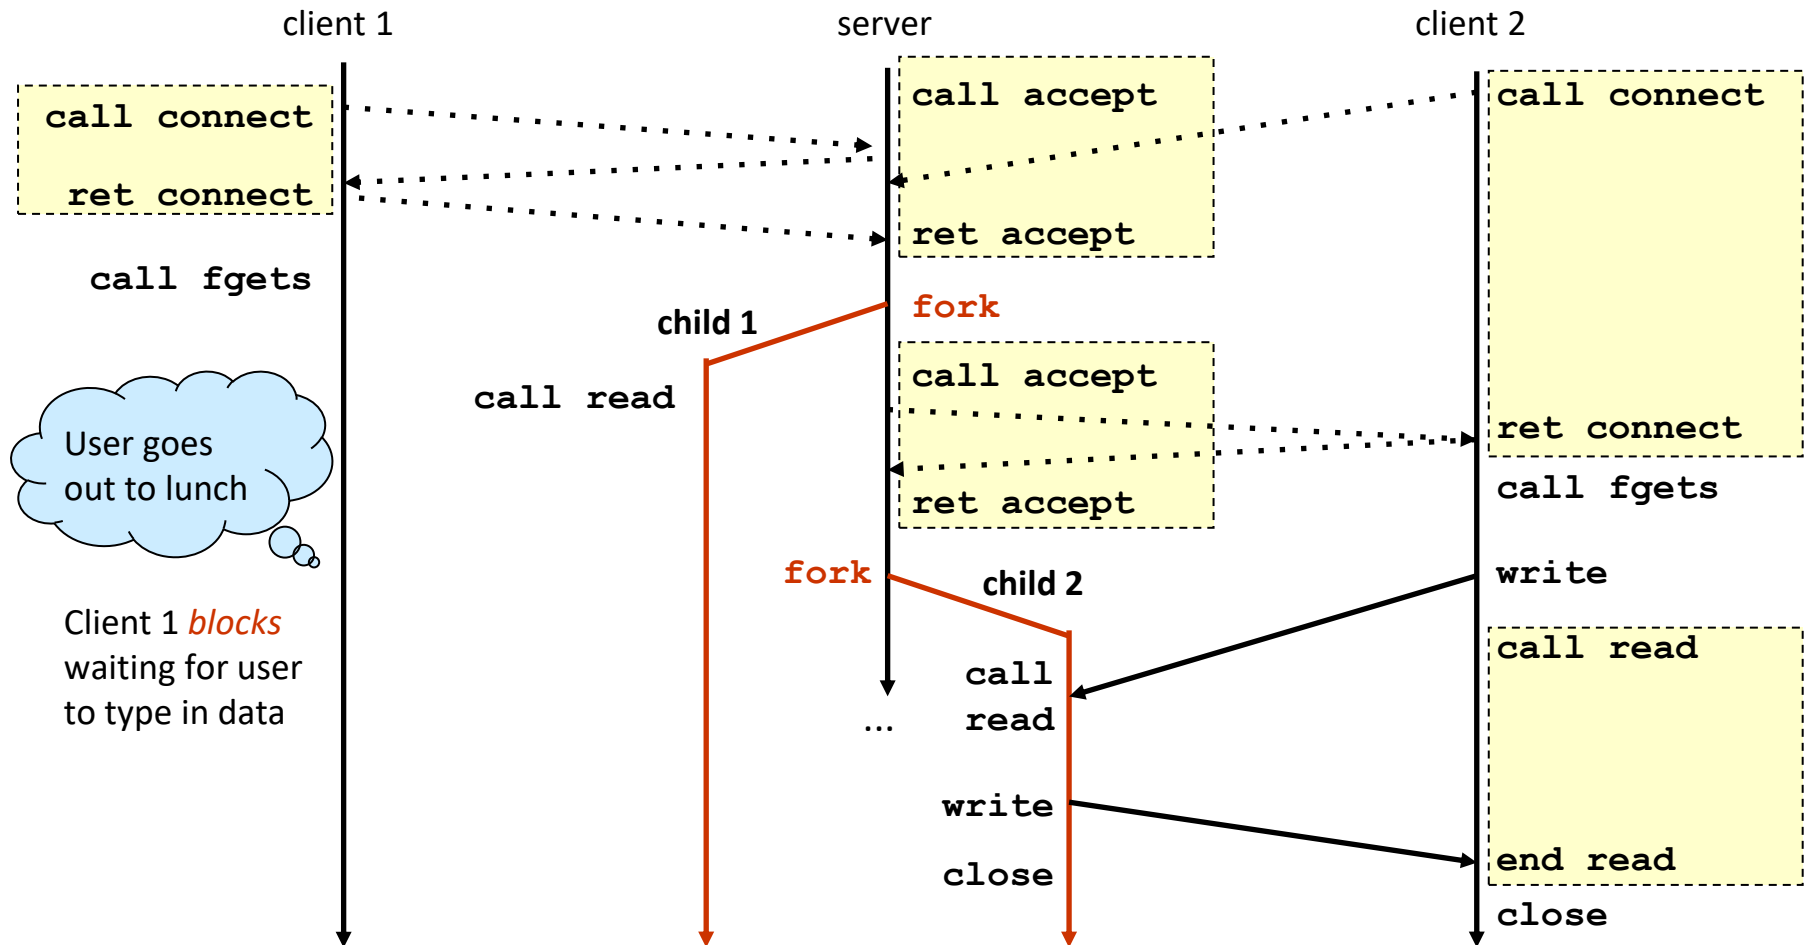
- Solution: use concurrent server
 - Use **multiple concurrent flows** to serve multiple clients at the same time

Creating Concurrent Flows

- Process-based
 - Kernel automatically interleaves multiple logical flows
 - Each flow has its own private address space
- Thread-based
 - Kernel automatically interleaves multiple logical flows
 - Each flow shares the same address space
- Event-based
 - User manually interleaves multiple logical flows
 - All flows share the same address space
 - Popular for high-performance server design
 - Using technique called **I/O multiplexing**

Approach #1: Process-based Server

- Concurrent server with multiple processes



Approach #1: Process-based Server

- Server must reap zombie children
 - To avoid fatal memory leak
- Server must close its copy of *connfd*
 - Kernel keeps reference for each socket/open file
 - After **fork()**, **refcnt(connfd) = 2**
 - Connection will not be closed until **refcnt(connfd) = 0**

Process-based Concurrent Echo Server

```
int main (int argc, char *argv[]) {  
    . . .  
    // parent process gets SIGCHLD, when child process terminates or stops  
    signal(SIGCHLD, handler);  
    while (1) {  
        if ((connfd = accept(listenfd, (struct sockaddr *)&caddr,  
                             (socklen_t *)&caddrlen)) < 0) {  
            printf("accept() failed\n");  
            continue;  
        }  
        if (fork() == 0) {  
            close(listenfd);  
            while ((n = read(connfd, buf, MAXLINE)) > 0) {  
                printf("got %d bytes from client.\n", n);  
                write(connfd, buf, n);  
            }  
            close(connfd);  
            exit(0);  
        }  
        close(connfd);  
    }  
}
```

```
void handler(int sig) {  
    pid_t pid;  
    int stat;  
    // wait for any child process that has changed state  
    while ((pid = waitpid(-1, &stat, WNOHANG)) > 0);  
    return;  
}
```


Pros and Cons of Process-based Design

■ Pros

- Handling multiple connections concurrently
- Clean sharing model
 - descriptors (no)
 - open file tables (yes)
 - global variables (no)
- Simple and straightforward

■ Cons

- Additional overhead for process control
 - Process creation, termination, and switching
- Nontrivial to share data between processes
 - Require IPC mechanisms

Approach #2: Thread-based Server

- Multiple threads within a process
 - Each thread shares same code, data, virtual address space of the process
 - Each thread has its own logical control flow, stack, thread ID (TID)
- vs. Process
 - Process does not share all code and data
 - Threads are somewhat less expensive than processes
 - Process control (creating and reaping) twice as expensive as thread control

Issues with Thread-based Server

- Must run “detached” to run independently and automatically free memory resources
 - **Joinable thread** can be reaped and killed by other threads
 - Must be reaped with `pthread_join()`
 - **Detached thread** cannot be reaped or killed by other threads
 - Resources are automatically reaped on termination
 - Exit state and return value are not saved
 - Default state is **joinable**
 - Use `pthread_detach(pthread_self())` to make **detached**

Issues with Thread-based Server

- Must be careful to avoid unintended sharing
 - e.g., what happens if we pass the address of **connfd** to the thread routine?

```
int connfd;
while (1) {
    connfd = accept ( ... );
    pthread_create(&tid, NULL, thread_func, &connfd);
}
```

- All functions called by a thread must be **thread-safe**
 - There is no problem in program execution
 - Even if more than one thread calls function at a time

Thread-based Concurrent Echo Server

```
int main (int argc, char *argv[]) {  
    . . .  
    int *connfdp;  
    pthread_id tid;  
    while (1) {  
        connfdp = (int *)malloc(sizeof(int));  
        if ((*connfdp = accept(listenfd, (struct sockaddr *)&caddr,  
                               (socklen_t *)&caddrlen)) < 0) {  
            printf("accept() failed\n");  
            free(connfdp);  
            continue;  
        }  
        pthread_create(&tid, NULL, thread_func, connfdp);  
    }  
}
```

```
void *thread_func(void *arg) {  
    int n;  
    char buf[MAXLINE];  
    int connfd = *((int*)arg);  
    pthread_detach(pthread_self());  
    free(arg);  
    while ((n = read(connfd, buf, MAXLINE)) > 0)  
        write(connfd, buf, n);  
    close(connfd);  
    return NULL;  
}
```

Pros and Cons of Thread-based Design

■ Pros

- Easy to share data structures between threads
- Threads are more efficient (cheaper) than processes

■ Cons

- Unintentional sharing can introduce subtle and hard-to-reproduce errors
 - Ease of data sharing is both the greatest strength and the greatest weakness of threads

Approach #3: Event-based Server

- Use I/O multiplexing technique
 - Provide more control with less overhead
- Event-based concurrent servers
 - Maintain pool of connected descriptors
 - Repeat following forever:
 - Use Unix **select()** system call to block until:
 - (a) New connection request arrives on the listening descriptor
 - (b) New data arrives on an existing connected descriptor
 - If (a), add new connection to the pool of connections
 - If (b), read any available data from the connection
 - Close connection on EOF and remove it from the pool

I/O Multiplexing Function

- `int select(int nfd, fd_set* readfds, fd_set* writefds, fd_set* exceptfds, struct timeval* timeout)`

- **nfd**

- The highest-numbered fd in any of three sets, plus 1

- **readfds / writefds / exceptfds**

- Sets of file descriptors are watched for reading/writing/exception

- **timeout**

- Interval that **select()** should block waiting for fds to become ready

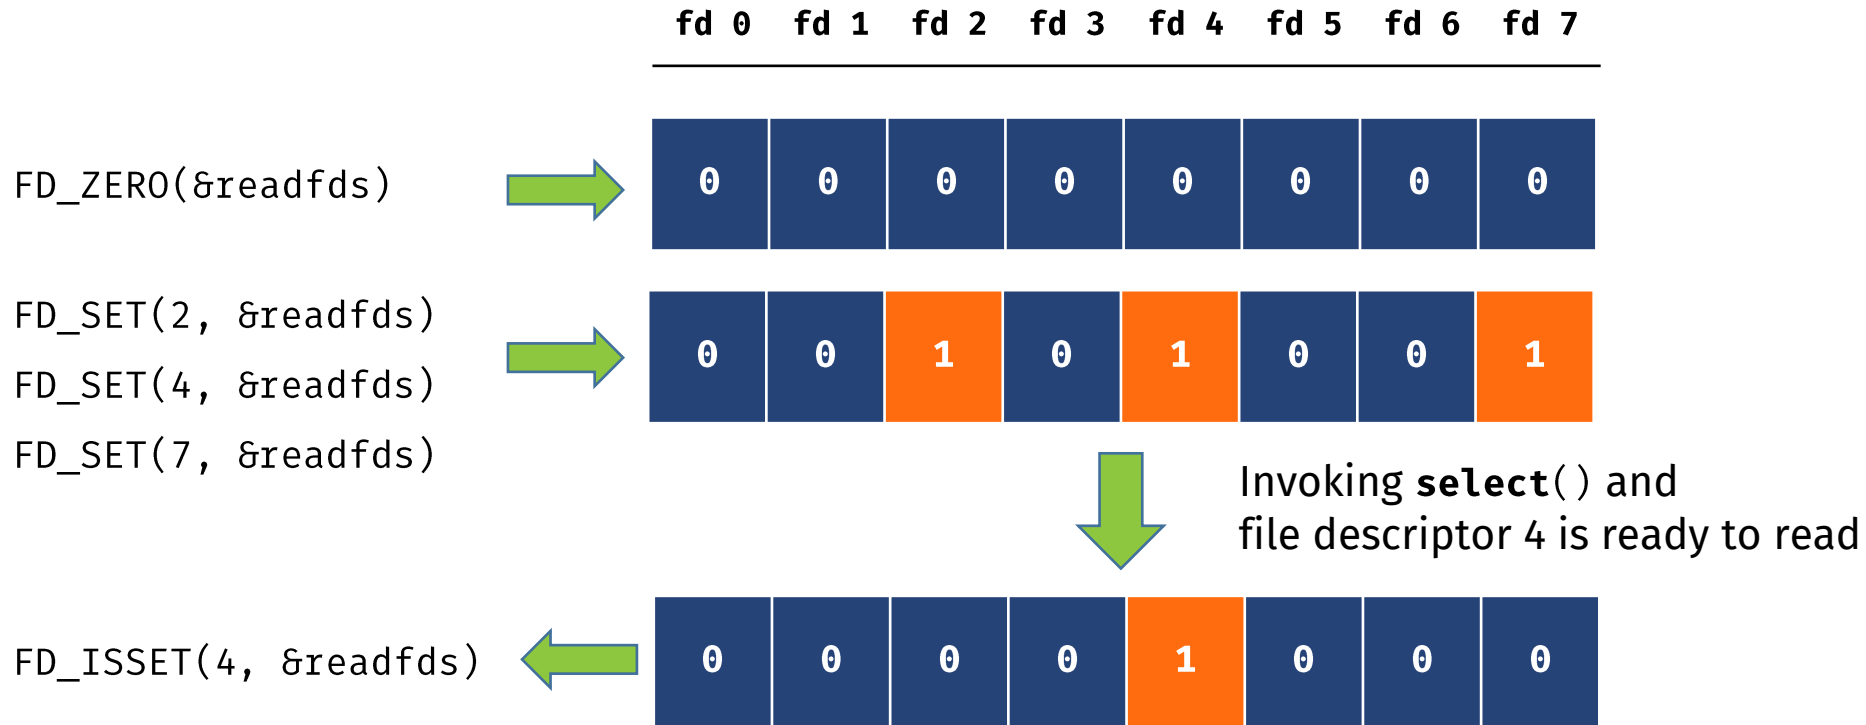
- Return value

- Return number of ready descriptors
 - And sets each bit of **readfds/writefds/exceptfds** to indicate the ready status of its corresponding descriptor

I/O Multiplexing Macros

- Macros for manipulating set descriptors
 - *void* **FD_ZERO**(*fd_set** **fdset**)
 - Turn off all bits in **fdset**
 - *void* **FD_SET**(*int* **fd**, *fd_set** **fdset**)
 - Turn on bit **fd** in **fdset**
 - *void* **FD_CLR**(*int* **fd**, *fd_set** **fdset**)
 - Turn off bit **fd** in **fdset**
 - *int* **FD_ISSET**(*int* **fd**, *fd_set** **fdset**)
 - Is bit **fd** in **fdset** turned on?
 - If on, return Non-zero
 - If off, return 0

I/O Multiplexing Macros Example



- *fd_set* does not remember the previous state
 - Value of previous *fd_set* must be stored before invoking **select()**

Event-based Concurrent Echo Server

```
#include <sys/select.h>

int main (int argc, char *argv[]) {
    . . .
    fd_set readset, copyset;
    FD_ZERO(&readset);
    FD_SET(listenfd, &readset);
    int fdmax = listenfd, fdnum;

    while (1) {
        copyset = readset;
        struct timeval timeout;
        timeout.tv_sec = 5;
        timeout.tv_usec = 0;

        if ((fdnum = select(fdmax + 1, &copyset, NULL, NULL, &timeout)) < 0) {
            printf("select() failed\n");
            exit(0);
        }

        if (fdnum == 0) {
            printf("Time out\n");
            continue;
        }
    }
}
```

Event-based Concurrent Echo Server

```
for (int i = 0; i < fdmax + 1; i++) {
    if (FD_ISSET(i, &copyset)) {
        if (i == listenfd) {
            if ((connfd = accept(listenfd, (struct sockaddr *)&caddr,
                                (socklen_t *)&caddrlen)) < 0) {
                printf ("accept() failed.\n");
                continue;
            }
            FD_SET(connfd, &readset);
            if (fdmax < connfd) fdmax = connfd;
        }
        else {
            if ((n = read(i, buf, MAXLINE)) > 0) {
                printf ("got %d bytes from client.\n", n);
                write(i, buf, n);
            }
            else {
                FD_CLR(i, &readset);
                printf("connection terminated.\n");
                close(i);
            }
        }
    }
}
```

Other I/O Multiplexing Function

- `int poll (struct pollfd *fds, nfds_t ndfs, int timeout)`
 - More efficient for large-valued or sparse file descriptors
 - **fds**
 - Set of file descriptors to be monitored is specified
 - Also contain requested events and returned events

```
struct pollfd {  
    int fd;      /* file descriptor */  
    short events; /* requested events */  
    short revents; /* returned events */  
}
```

- **ndfs**
 - Specify number of items in **fds** array
- **timeout**
 - Interval that **poll()** should block waiting for **fds** to become ready
- Return value
 - Number of structures which have nonzero revents fields

Other I/O Multiplexing Function

- `int epoll_wait (int epfd, struct epoll_event *events, int maxevents, int timeout)`
 - O(1) operation
 - `epoll_create()`, `epoll_ctl()`, and wait for events using `epoll_wait()`
 - **epfd**
 - epoll file descriptor (return value from `epoll_create()`)
 - **events**
 - Return address of `epoll_event` structures that contain ready fd
 - **maxevents**
 - Specify that maximum number of events to process at a time
 - **timeout**
 - Interval that `epoll_wait()` should block waiting for fd to become ready

Pros and Cons of Event-based Server

■ Pros

- One logical control flow and address space
- Can single-step with a debugger
- No process or thread control overhead
 - Design of choice for high-performance servers

■ Cons

- Significantly more complex to code than process- or thread-based designs
- Can be vulnerable if malicious client sends particular text and halts

Lab Exercise

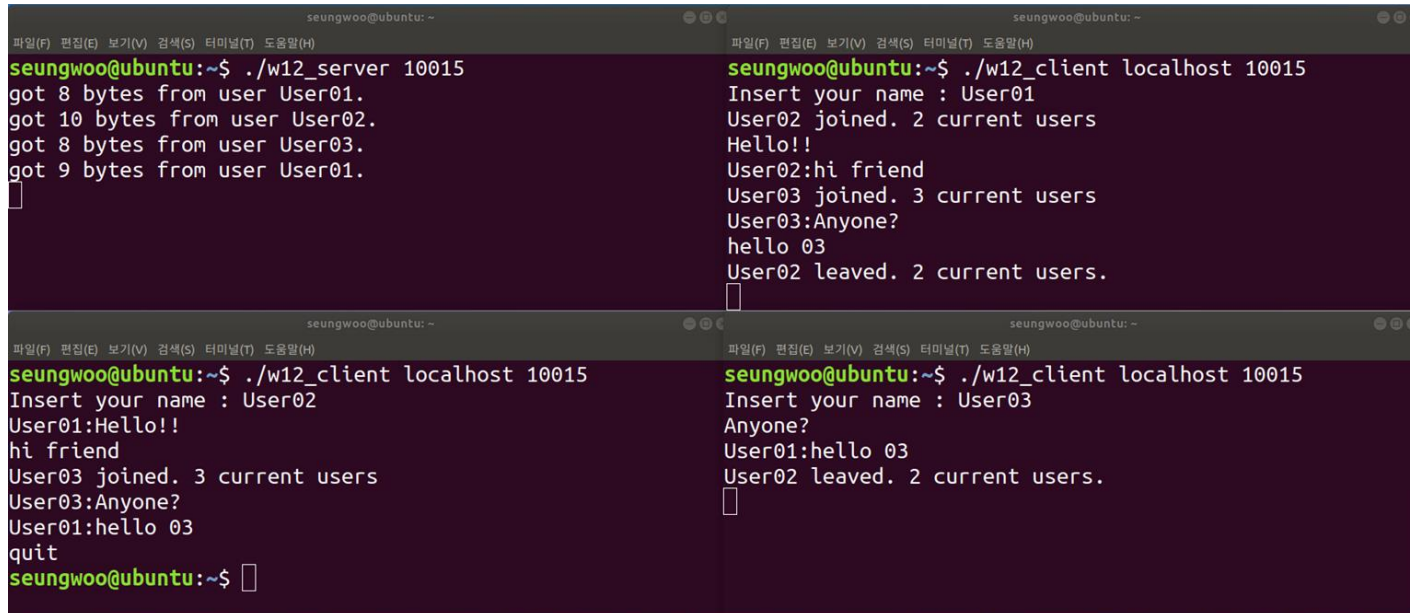
- Make *chatting room* (Up to 10 users) using `select()`
 - Client should insert the **name** (Up to 10 bytes) of itself
 - Client should be under an infinite loop with `conditional exit` (“quit”)
 - When client sends message (Up to 80 bytes),
server receives it and sends it to the other clients
 - When client joins or leaves chatting room, server sends message with number of current clients to other clients
 - Use `FD_ISSET()` to check if there is input in specific file descriptor
 - **No timeout** is needed
 - e.g., `select(nfds, &readset, NULL, NULL, NULL);`

Lab Exercise

- Skeleton code

- `cp ~swe2024-41_23s/2023s/p12 ./ -r`

- Example)



```
seungwoo@ubuntu: ~  
파일(F) 편집(E) 보기(V) 검색(S) 터미널(T) 도움말(H)  
seungwoo@ubuntu:~$ ./w12_server 10015  
got 8 bytes from user User01.  
got 10 bytes from user User02.  
got 8 bytes from user User03.  
got 9 bytes from user User01.  
█
```

```
seungwoo@ubuntu: ~  
파일(F) 편집(E) 보기(V) 검색(S) 터미널(T) 도움말(H)  
seungwoo@ubuntu:~$ ./w12_client localhost 10015  
Insert your name : User01  
User02 joined. 2 current users  
Hello!!  
User02:hi friend  
User03 joined. 3 current users  
User03:Anyone?  
hello 03  
User02 leaved. 2 current users.  
█
```

```
seungwoo@ubuntu: ~  
파일(F) 편집(E) 보기(V) 검색(S) 터미널(T) 도움말(H)  
seungwoo@ubuntu:~$ ./w12_client localhost 10015  
Insert your name : User02  
User01:Hello!!  
hi friend  
User03 joined. 3 current users  
User03:Anyone?  
User01:hello 03  
quit  
seungwoo@ubuntu:~$ █
```

```
seungwoo@ubuntu: ~  
파일(F) 편집(E) 보기(V) 검색(S) 터미널(T) 도움말(H)  
seungwoo@ubuntu:~$ ./w12_client localhost 10015  
Insert your name : User03  
Anyone?  
User01:hello 03  
User02 leaved. 2 current users.  
█
```

Exercise Submission

- Submit your exercise code and Makefile
 - InUiYeJi cluster
 - Submit the folder into p12
 - ~swe2024-41_23s/bin/submit p12 p12
 - Due date: Tonight, 17 May 2023, 23:59
 - We will compile by using command *make*
 - If compilation fails, your points for this exercise will be **zero**

```
./p12
```

```
server.c  
client.c  
Makefile
```

Summary Report

- Summary report about man command result of
 - **select()**
 - **poll()**
 - **epoll_wait()**
- Submission form
 - A4 size PDF format (No page limitation)
 - [SWE2024 Report-10] studentID_name
 - Ex) [SWE2024 Report-10] 2022XXXXXX_홍길동
 - Submit to iCampus
 - Due by Friday, 19 May 2023, 23:59