

Synchronization I

Prof. Joonwon Lee (joonwon@skku.edu)

TA – Jaehyung Park (jaeseanpark@gmail.com)

TA – Luke Albano (lukealbano@arcs.skku.edu)

Sungkyunkwan University

Warm Up Example

```
#include <stdio.h>
#include <pthread.h>

int num;

void *inc (void *tid) {
    int iter = 10000;

    while(iter--) num++;
}

void *dec (void *tid) {
    int iter = 10000;

    while(iter--) num--;
}
```

```
int main() {
    int t;

    for(t = 0; t < 10; t++) {
        num = 0;
        pthread_t thread_inc, thread_dec;
        pthread_create(&thread_inc, NULL, &inc, NULL);
        pthread_create(&thread_dec, NULL, &dec, NULL);

        pthread_join(thread_inc, NULL);
        pthread_join(thread_dec, NULL);

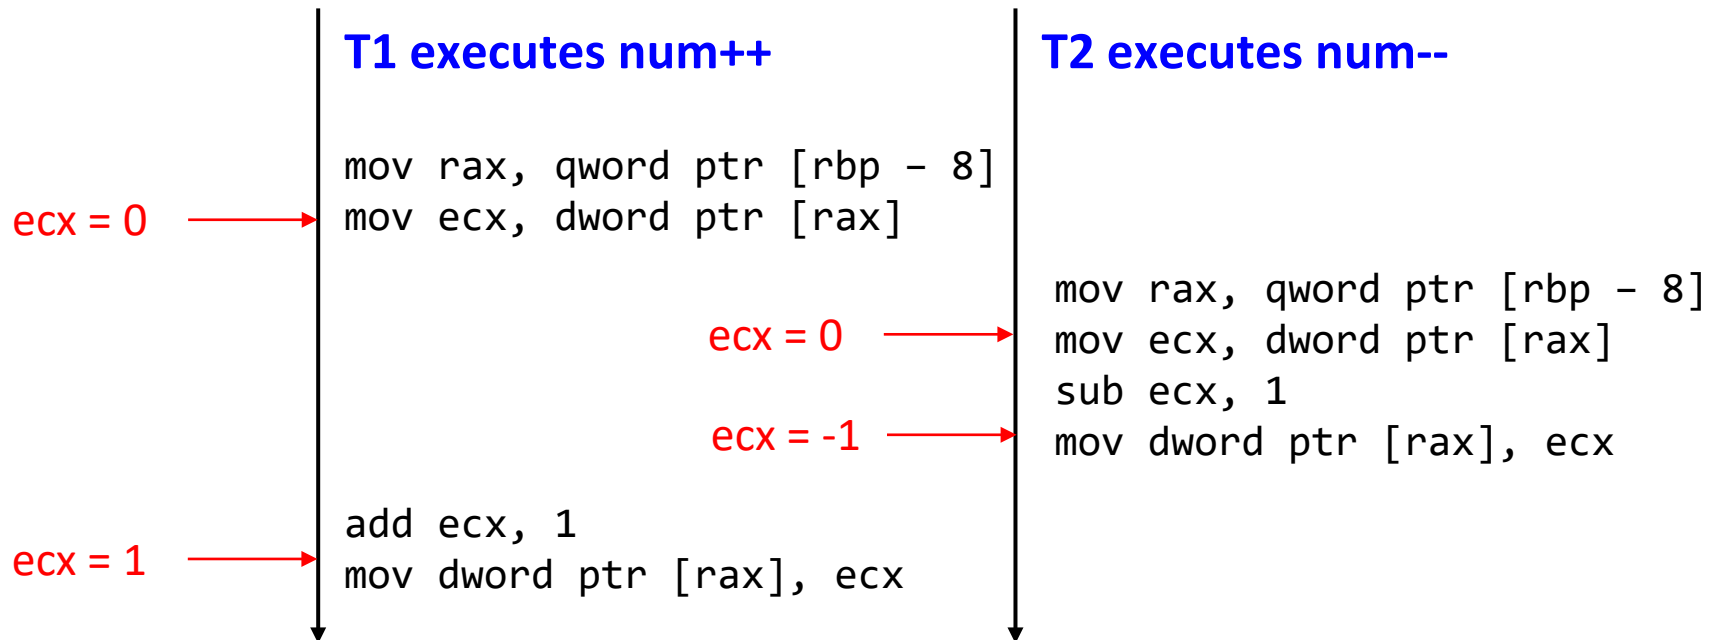
        printf("#%d, %d\n", t, num);
    }
    return 0;
}
```

Warm Up Example

```
#0, -2759  
#1, 1200  
#2, -10000  
#3, 0  
#4, 10000  
#5, -264  
#6, -584  
#7, -662  
#8, -5587  
#9, -3864
```

Instruction-level Timeline

num = 0



Mutex

- Mutex is abbreviation for "mutual exclusion"
 - Primary means of implementing **thread synchronization**
 - Protects shared data when **multiple writes** occurs
 - A mutex variable acts like "lock" protecting access **to shared resource**
 - Only one thread can lock (or own) mutex variable at any given time
 - Even if several threads try to lock mutex, only one thread will be successful
 - Other threads are blocked until owner releases mutex
 - Mutex is used to prevent "race" conditions
 - Race condition: anomalous behavior due to unexpected critical dependence on the relative timing of events

Mutex Creation / Termination API

- Static initialization

- `pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;`

- Dynamic initialization

- `pthread_mutex_t m;`

- `pthread_mutex_init(&m, (pthread_mutexattr_t*)NULL);`

- Mutex termination

- `pthread_mutex_destroy(&m);`

- Destroy mutex object

- Free resources it might hold

Mutex Operation APIs

- `int pthread_mutex_lock (pthread_mutex_t *m)`
 - **Acquire lock** on specified mutex variable
 - If the **mutex** is already locked by another thread, block calling thread until **mutex** is unlocked
- `int pthread_mutex_unlock (pthread_mutex_t *m)`
 - Unlock **mutex** if called by owning thread
- `int pthread_mutex_trylock (pthread_mutex_t *m)`
 - Attempt to lock **mutex**
 - If the **mutex** is already locked, **return immediately with "busy" error code**

Mutex Example

```
int num;
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;

void *inc (void *tid) {
    int iter = 10000;

    while(iter--) {
        pthread_mutex_lock(&m);
        num++;
        pthread_mutex_unlock(&m);
    }
}

void *dec (void *tid) {
    int iter = 10000;

    while(iter--) {
        pthread_mutex_lock(&m);
        num--;
        pthread_mutex_unlock(&m);
    }
}
```

```
#1 0
#2 0
#3 0
#4 0
#5 0
#6 0
#7 0
#8 0
#9 0
#10 0
```


Condition Variables (CV)

- Another way for thread synchronization
 - While mutexes implement synchronization by controlling thread access to data, condition variables allow threads to synchronize based upon actual value of data
 - Without condition variables, programmer would need to have threads continually polling to check if condition is met
 - This can be very resource consuming since thread would be continuously busy in this activity
 - A condition variable is always used in conjunction with mutex lock

Operation Mechanism of CV

- How condition variables work
 - Thread locks mutex associated with condition variable
 - Thread tests condition to see if it can proceed
 - If it can
 - Your thread does its work
 - Your thread unlocks mutex
 - If it cannot
 - The thread sleeps & **mutex is automatically released**
 - Some other threads signal condition variable
 - Your thread wakes up from waiting **with mutex automatically locked**, and it does its work
 - Your thread releases mutex when it is done

CV Creation / Termination API

- Static initialization

- `pthread_cond_t cond = PTHREAD_COND_INITIALIZER;`

- Dynamic initialization

- `pthread_cond_t cond;`

- `pthread_cond_init(&cond, (pthread_condattr_t*)NULL);`

- Condition variable termination

- `pthread_cond_destroy(&cond);`

- Destroy condition variable object

- Free resources it might hold

CV Operation APIs

- `int pthread_cond_wait (pthread_cond_t *cond, pthread_mutex_t *mutex)`
 - Block calling thread until specified condition is signaled
 - This should be called while mutex is locked, and it will automatically release mutex while it waits
- `int pthread_cond_signal (pthread_cond_t *cond)`
 - Signal another thread which is waiting on CV
 - Calling thread should have a lock
- `int pthread_cond_broadcast (pthread_cond_t *cond)`
 - Used if more than one thread is in blocking wait state

Condition Variable Example

```
action() {  
    pthread_mutex_lock(&mutex);  
    while(x != 0)  
        pthread_cond_wait(&cond, &mutex);  
    real_action();  
    pthread_mutex_unlock(&mutex);  
}
```

```
counter() {  
    pthread_mutex_lock(&mutex);  
    x--;  
    if(x == 0)  
        pthread_cond_signal(&cond);  
    pthread_mutex_unlock(&mutex);  
}
```

Thread Safety

- Thread-safe
 - Functions called from thread must be **thread-safe**
 - We identify four (non-disjoint) classes of thread-unsafe functions:
 - Class 1: Failing to protect shared variables
 - Class 2: Relying on persistent state access invocations
 - Class 3: Returning pointer to static variable
 - Class 4: Calling thread-unsafe functions

Thread Safety Class #1

- Failing to protect shared variables
 - Fix: use mutex operations
 - Issue: synchronization operations will slow down code

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
int cnt = 0;

/* Thread routine */
void *count(void *arg) {
    int i;

    for (i=0; i<NITERS; i++) {
        pthread_mutex_lock (&lock);
        cnt++;
        pthread_mutex_unlock (&lock);
    }
    return NULL;
}
```

Thread Safety Class #2

- Relying on persistent state across multiple function invocations
 - Random number generator relies on static state
 - Fix: rewrite function so that caller passes in all necessary state

```
/* rand - return pseudo-random integer on  
0..32767 */  
static unsigned int next = 1;
```

```
int rand(void) {  
    next = next * 1103515245 + 12345;  
    return (unsigned int)(next/65536) %  
32768;  
}
```

```
/* srand - set seed for rand() */  
void srand(unsigned int seed) {  
    next = seed;  
}
```

```
/* rand - return pseudo-random integer on  
0..32767 */
```

```
int rand_r(int *nextp) {  
    *nextp = *nextp * 1103515245 + 12345;  
    return (unsigned int)(*nextp/65536) %  
32768;  
}
```



Thread Safety Class #3

■ Returning pointer to static variable

– Case #1

- Fix: rewrite code so caller passes pointer to **struct**
- Issue: require changes in caller and callee

```
struct hostent
*gethostbyname(char *name) {
    static struct hostent h;
    <contact DNS and fill in h>
    return &h;
}
```

```
hostp = malloc(...));
gethostbyname_r(name, hostp);
```

– Case #2

- Fix: lock-and-copy
- Issue: require only simple changes in caller (and none in callee)
 - ✓ However, caller must free memory

```
struct hostent
*gethostbyname_ts(char *name) {
    struct hostent *unshared
        = malloc(...);
    pthread_mutex_lock(&lock); /* lock */
    shared = gethostbyname(name);
    *unshared = *shared; /* copy */
    pthread_mutex_unlock(&lock);
    return unshared;
}
```

Thread Safety Class #4

- Calling thread-unsafe functions
 - Calling one thread-unsafe function makes entire function thread-unsafe
 - Fix: modify function so it calls only thread-safe functions

Exercise


- Monte Carlo method
- Single thread version

```
count = 0
r = 1

while i = 1~point_total
    x = random number between 0 and 1
    y = random number between 0 and 1

    if (x,y) within circle
        count += 1

pi = 4 * count / point_total
```



Exercise

- The parallel version gets two command line arguments
 - The number of threads
 - The number of random points per thread
- Main thread collects results from created threads and calculates value for pi
 - Each created thread must create random points and count the number of points that are inside circle
- Copy the skeleton code to your directory
`$ cp ~swe2024-41_23s/2023s/p13_skeleton.c ./`

```
> ./p13 10 10000000  
pi: 3.142123
```

Exercise

- Submit your exercise source code
 - To InUiYeJi Cluster
 - Put your Makefile and *.c files in p13 folder
 - Submit using

```
$ ~swe2024-41_23s/bin/submit p13 p13
```
 - We will compile by using command *make*
 - When compilation fails, you get zero points
 - Compiled binary name should be “*p13*”
- Due 2023/5/24 23:59